

ECE437/CS481

M02A: PROCESSES & THREADS

PROCESSES CONCEPTS

Chapter 3.1-3.2

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a small upward curve on the left, dipping into a V-shape in the center, and then curving back up on the right before continuing as a straight line to the edge.

□ What is a **program**?

- “A program contains **codes & static data stored in a file**” (von Neumann architecture)
 - ✓ Source file/code - before compilation
 - ✓ Binary file/code - ready to be loaded

□ What is a **process**?

- “A process is **a program that is being executed**”
 - ✓ A program can execute many times, so corresponding to many processes
 - ✓ Different process have their own **address spaces & process contexts**

Process Context

□ Process context includes

➤ Execution information

- ✓ **CPU register set image** (e.g., Stack Pointer (SP), Instruction Pointer (IP), etc.)
- ✓ **Process stack** - containing temporary data (e.g., subroutine arguments, temporary variables, return addresses)

%rax	Return value
%rbx	Callee saved
%rcx	Argument #4
%rdx	Argument #3
%rsi	Argument #2
%rdi	Argument #1
%rsp	Stack pointer
%rbp	Callee saved

%r8	Argument #5
%r9	Argument #6
%r10	Caller saved
%r11	Caller Saved
%r12	<u>Callee saved</u>
%r13	Callee saved
%r14	Callee saved
%r15	Callee saved

Process Context

□ Process context includes

➤ Environment information

- ✓ Memory address space or memory map, i.e., regions of memory that have been allocated to the process
- ✓ Resources, e.g., open file table

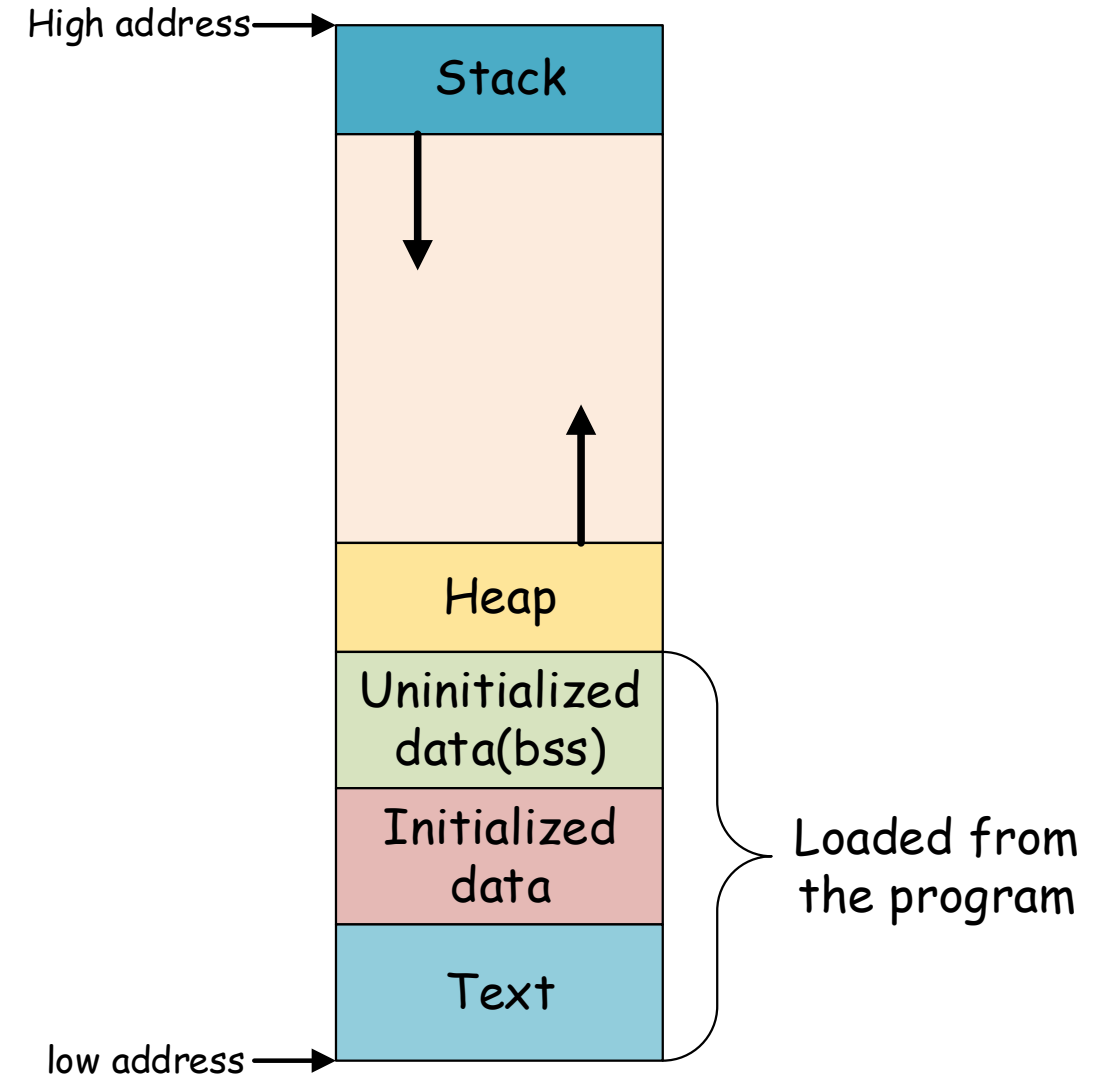
➤ System attributes

- ✓ process identification or Process ID (PID)
- ✓ process state information, e.g., waiting, ready, etc.
- ✓ process control information, e.g., priority

Process Memory Map

□ Process's **memory layout** in Unix/Linux

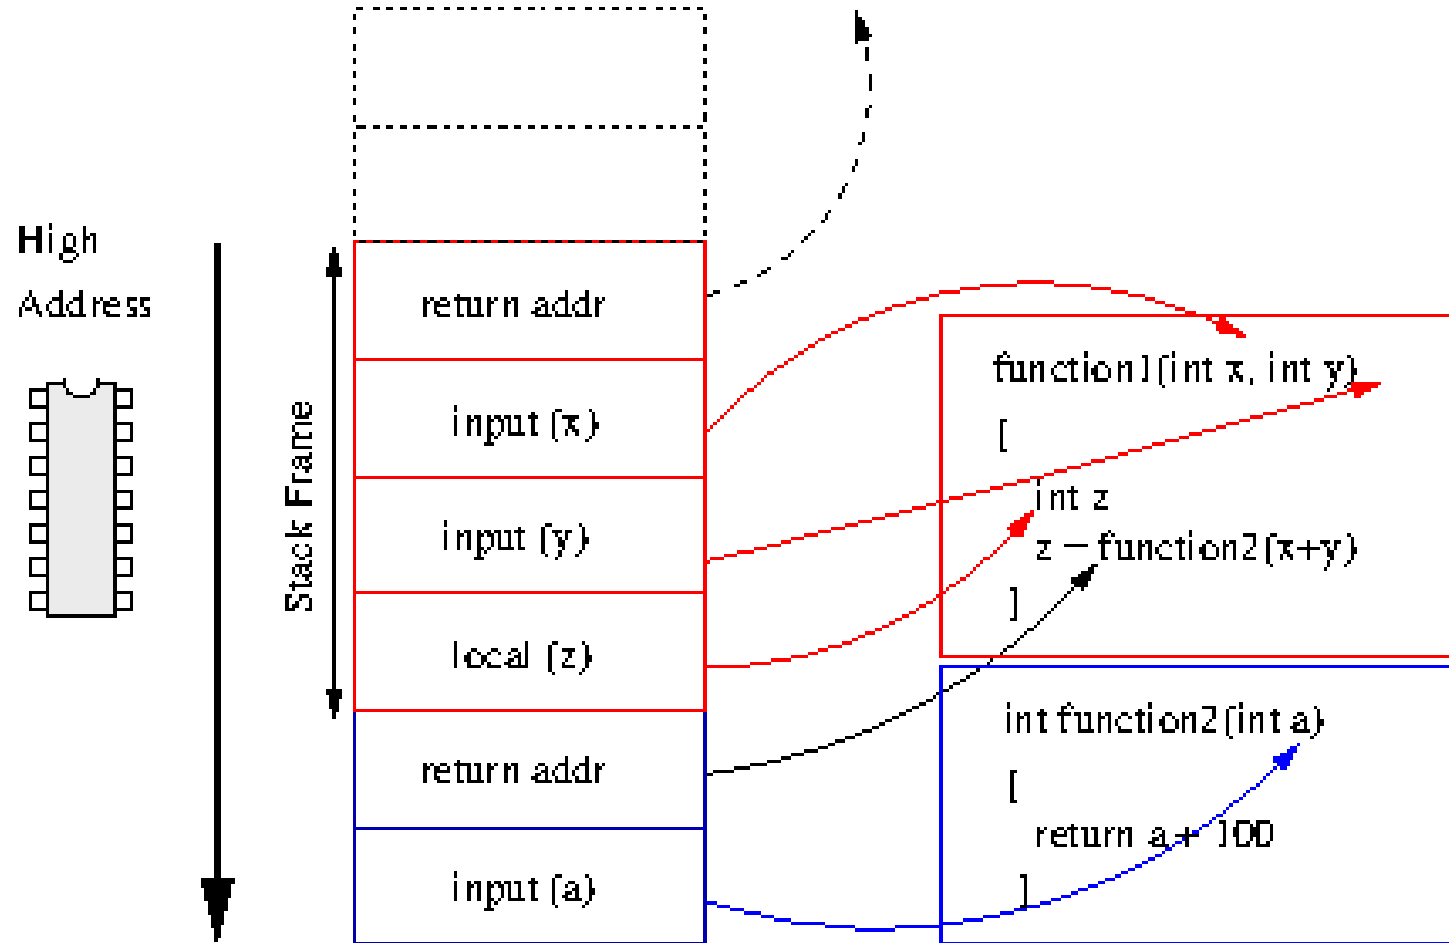
- **Text (code)**: the machine instructions to be executed. Usually, it stays unmodified over the lifecycle of the process.
- **Initialized data**
- **Uninitialized data**
- **Heap**: dynamically allocated memory
 - ✓ grow via a system call to request more memory
- **Stack**: memory space to manage function calls, returns, parameter passing, and local variables.
 - ✓ grow and shrink as the depth of function calls varies



Process Memory Map

□ Process's **memory layout** in Unix/Linux

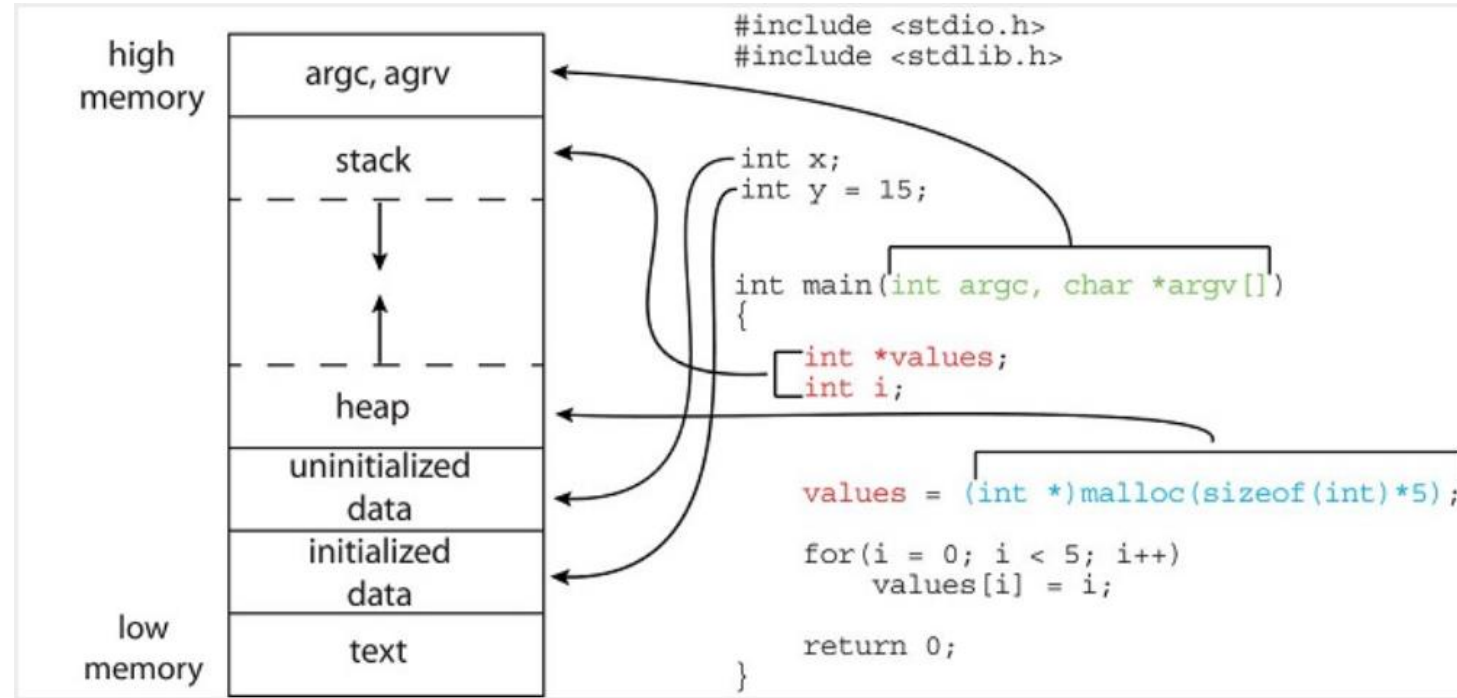
- Once a subfunction is invoked, a new **stack frame** is created.
- The stack frame is an area of memory which usually contains the return address, input arguments to the subfunction, and local variables.
- The stack starts at a high address in memory and progressively gets lower.



Process Memory Map

□ An example of memory layout of a C program

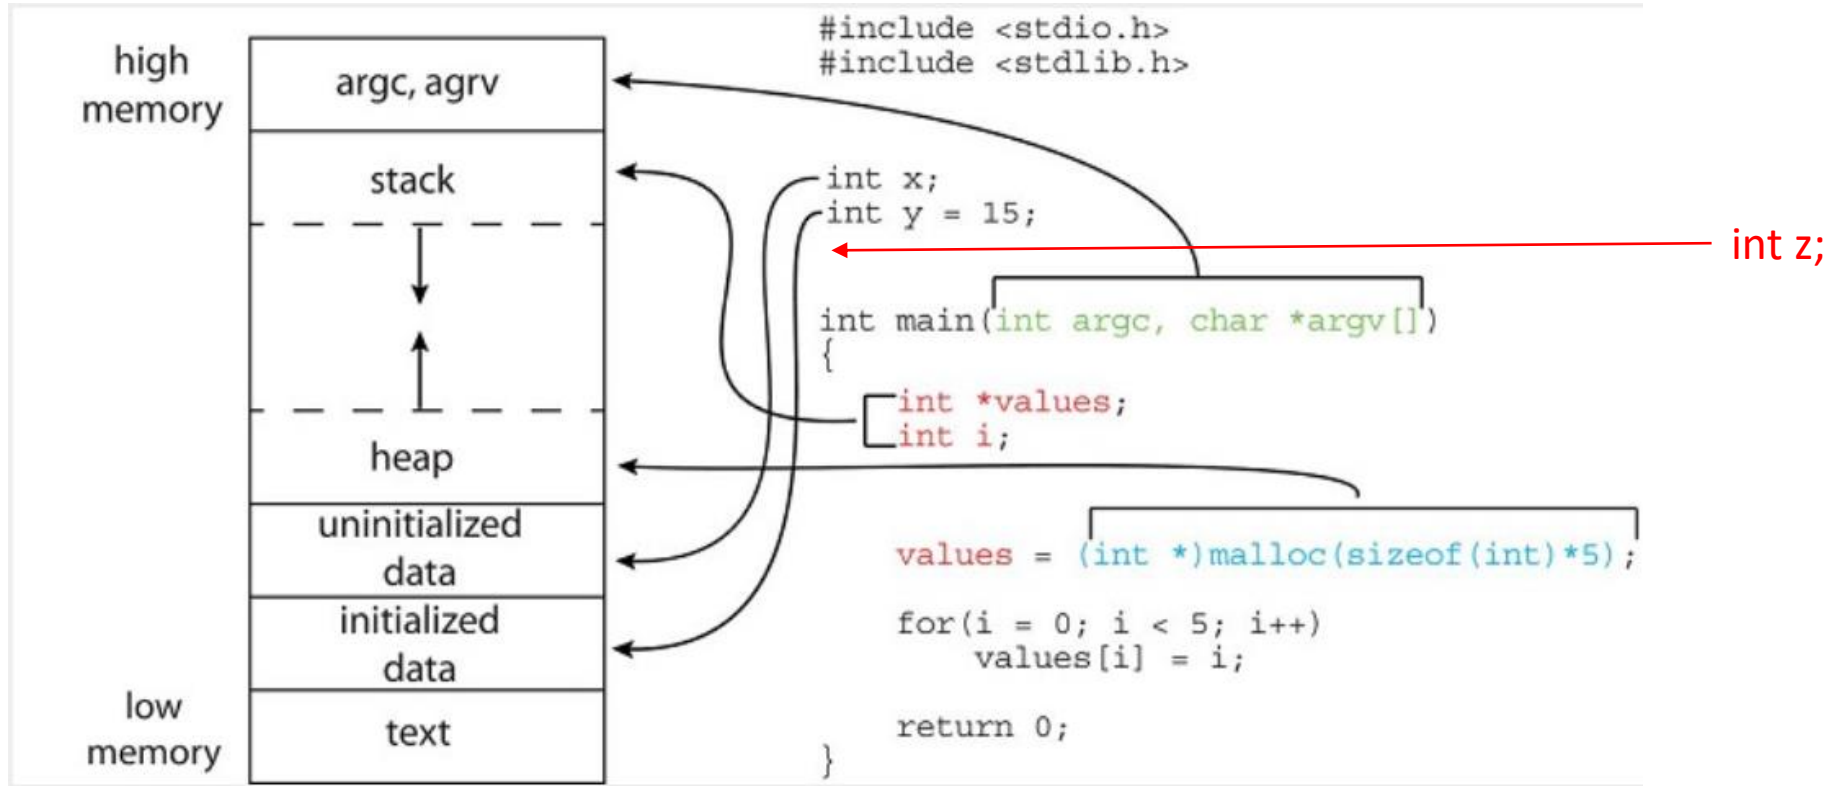
- The global data section is divided into different sections for (a) initialized data and (b) uninitialized data.
- The arguments and local variables in the main() function are stored in the stack.
- Malloc() create a memory space for heap.
- Run `size <fileName>` can see size (in bytes) of different parts.



text	data	bss	dec	hex	filename
1158	284	8	1450	5aa	memory

Process Memory Map

□ An example of memory layout of a C program



text	data	bss	dec	hex	filename
1158	284	8	1450	5aa	memory

Process State

❑ New (Start)

- The process is being created but has **not yet admitted** to the pool of executable processes
 - ✓ Admit --- go to Ready state: The OS will move a process to the Ready state when the system has enough memory space.

❑ Ready

- The process is waiting to be assigned to a processor.

❑ Terminated (Exit):

- The process has finished the execution.

❑ Running

- The process is currently being executed.
- In a single processor computer, at most one process can be in this state. But it is not true for today's multicore systems.

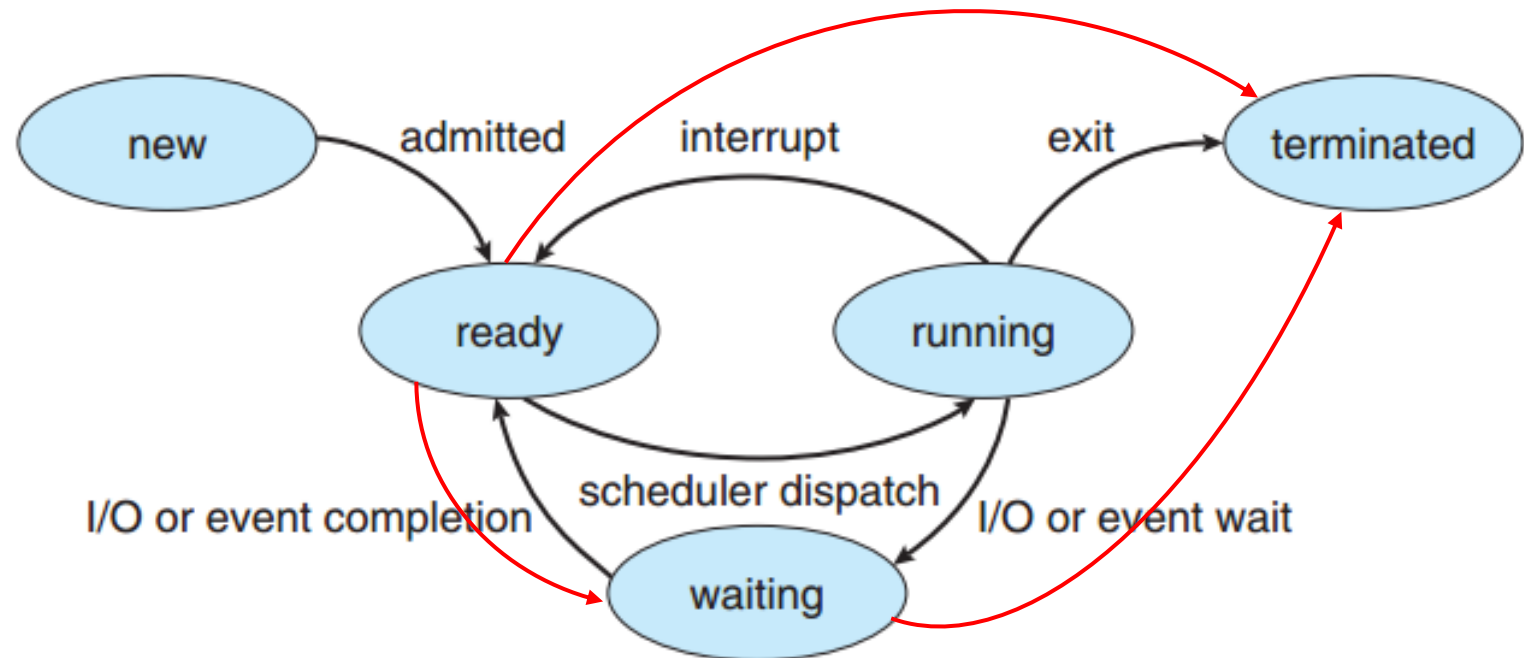
❑ Waiting (Blocked)

- The process is waiting for some event to occur.

Process State

□ Ready→Running/Waiting/Terminated

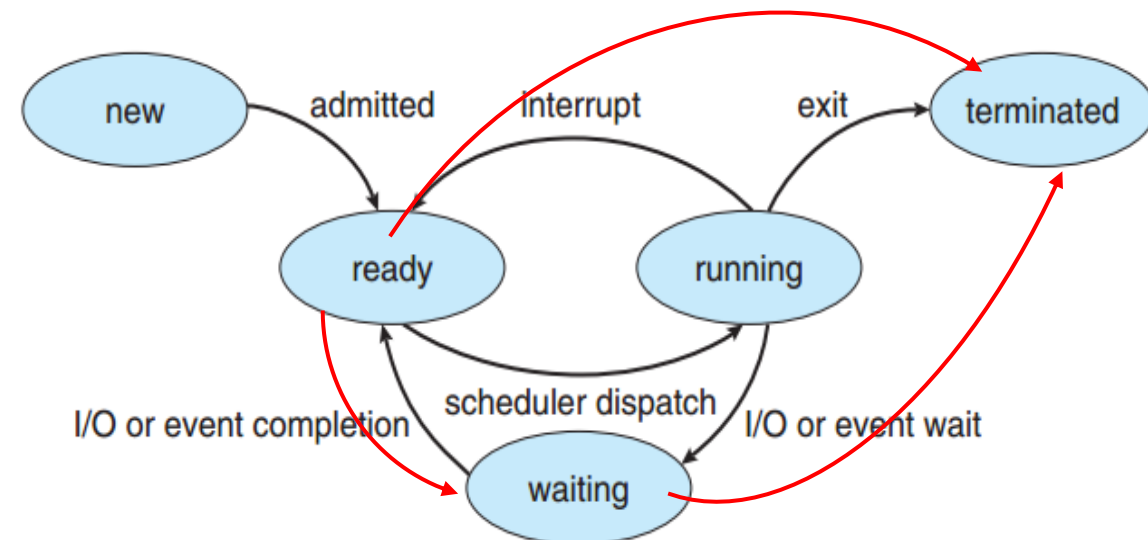
- Scheduler dispatch --- go to the Running state
- Suspend --- go to the Waiting state
- Kill --- go to the Terminated state



Process State

❑ Running → Ready/Waiting/Terminated


- Yield (i.e., the running process voluntarily release control of the processor) --- go to the Ready state
- Time slice is up (i.e., the running process has reached the maximum allowable time) --
- go to the Ready state
- Arrival of high priority process (i.e., the running process will be preempted due to its lower priority) --- go to the Ready state
- I/O request --- go to the Waiting state
- Suspend --- go to the Waiting state
- Terminate --- go to the Terminated state



Process State

□ Waiting→Ready/Terminated

- I/O complete, wakeup --- go to the Ready state
- Kill --- go to the terminated state



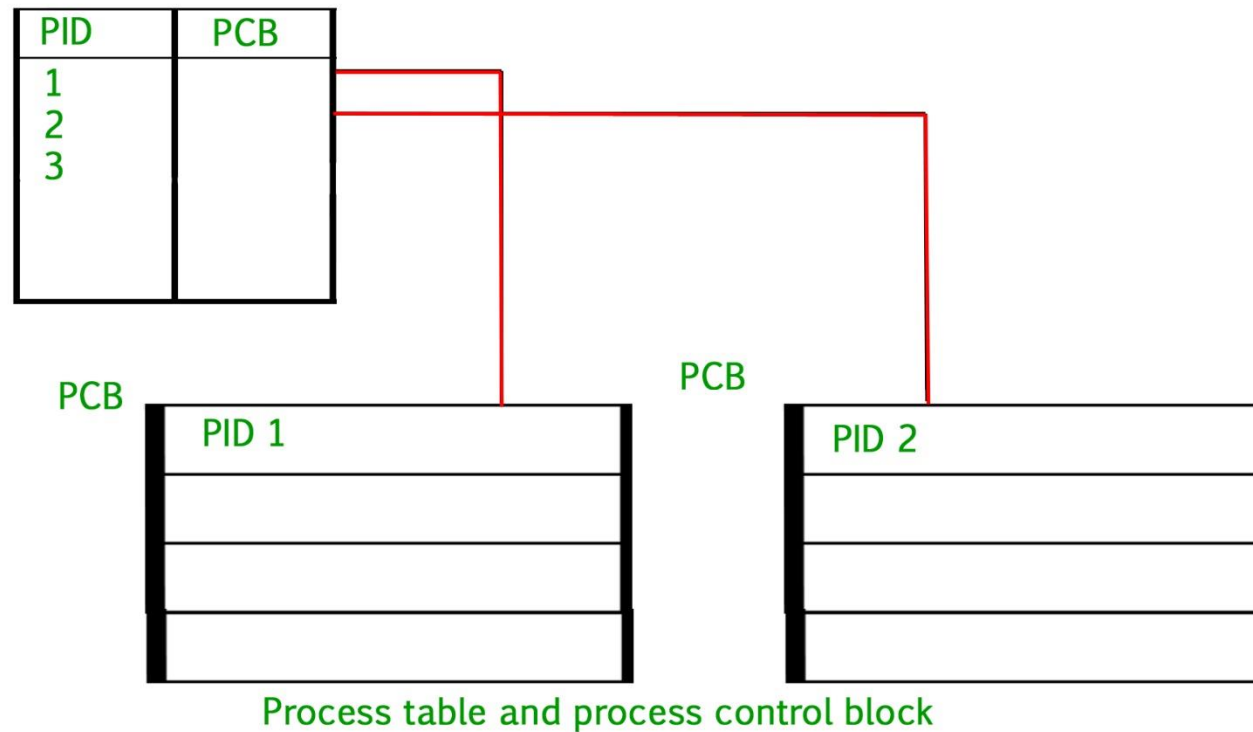
❑ Process Control Block (PCB)

- Each process is represented in the OS by a PCB
- An important data structure to keep the context of a process
 - ✓ **PID**: a unique integer as process ID
 - ✓ **UID**: user ID who is executing this process
 - ✓ **Program Counter**: A pointer to the address of the next instruction to be executed for this process.
 - ✓ **Process state**: the current process state
 - ✓ **Event**: any event for which the process may be waiting
 - ✓ **Memory management info**: information of page table, memory limits, segment table, etc.
 - ✓ **CPU registers**: values of CPU registers where a process is swapped out.
 - ✓ **Scheduling priority** : parameters for the process scheduling
 - ✓ **Accounting info**: timing, usage, ...
 - ✓ **IO status information**: a list of I/O devices allocated to the process.
- In linux:
 - ✓ defined in task_struct (/usr/src/linux/include/linux/sched.h)---over 95 fields!!!

Process State

❑ Process table/list

- The OS maintains a process table/list, which contains PIDs and pointers.

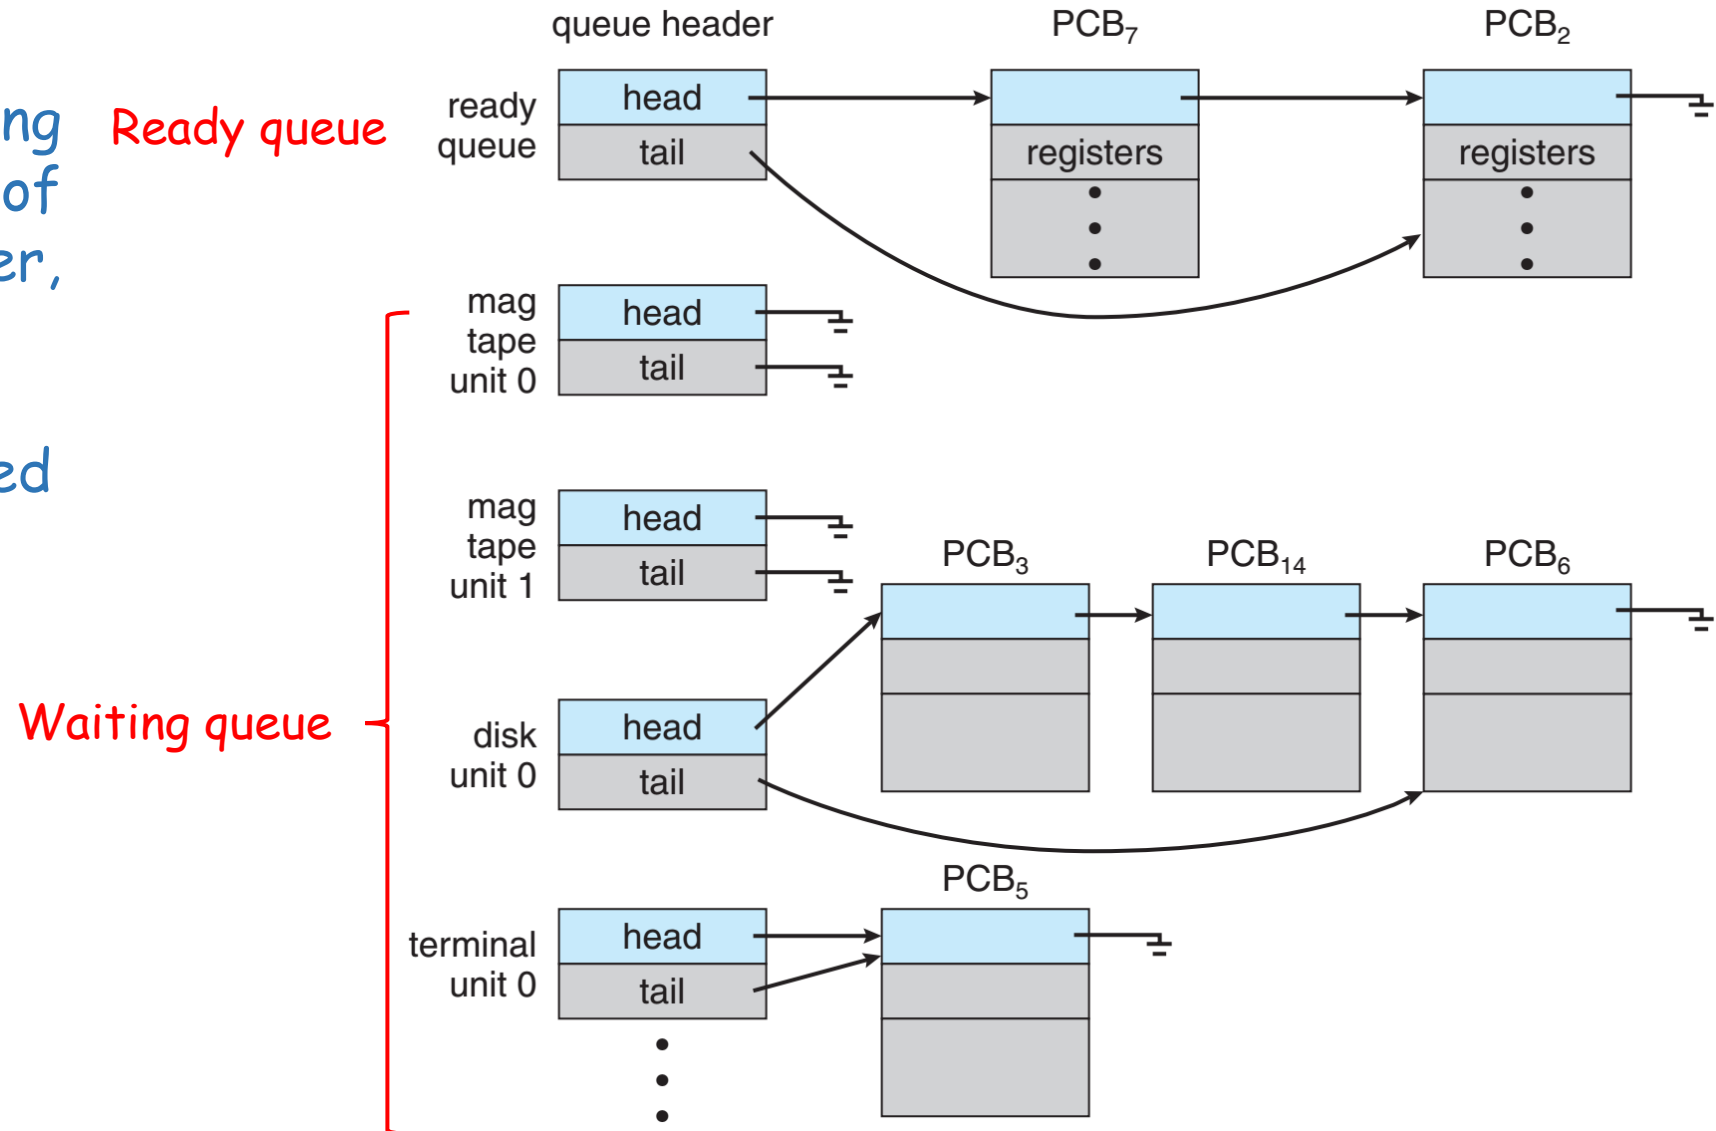


State Queues

- ❑ The OS maintains a collection of queues that represent the state of all processes in the system
 - Typically, one queue stands for a state, e.g., ready,
 - Each PCB is queued onto a specific state queue based on its status.
 - Once a process changes its state, its PCB is unlinked from one queue, and linked onto another.

State Queues

- There may be many waiting queues, one for each type of wait (particular device, timer, message, ...)
- Running queue? Terminated queue?



Process Context Switch

□ Process context switch

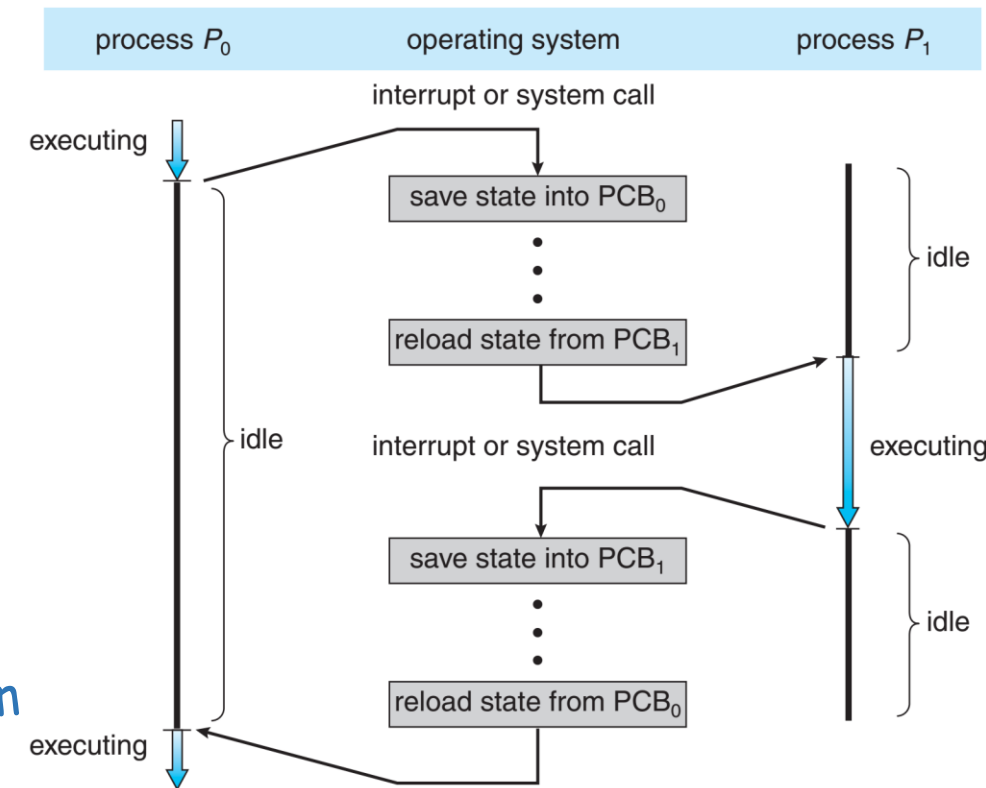
---Switching of the CPU from process A to another, requiring storing the state of process A such that it can be restored, and execution resumed from the same point later.

➤ Why need context switch

--- improve CPU utilization

➤ Steps of context switch

1. save the registers' data of the current process.
2. change process state (e.g., running->waiting) and put it to corresponding queue.
3. select another process from the ready queue.
4. change the newly selected process' state (e.g., ready->running) and reload the registers' data from its PCB.
5. (re)start computation of the new process.



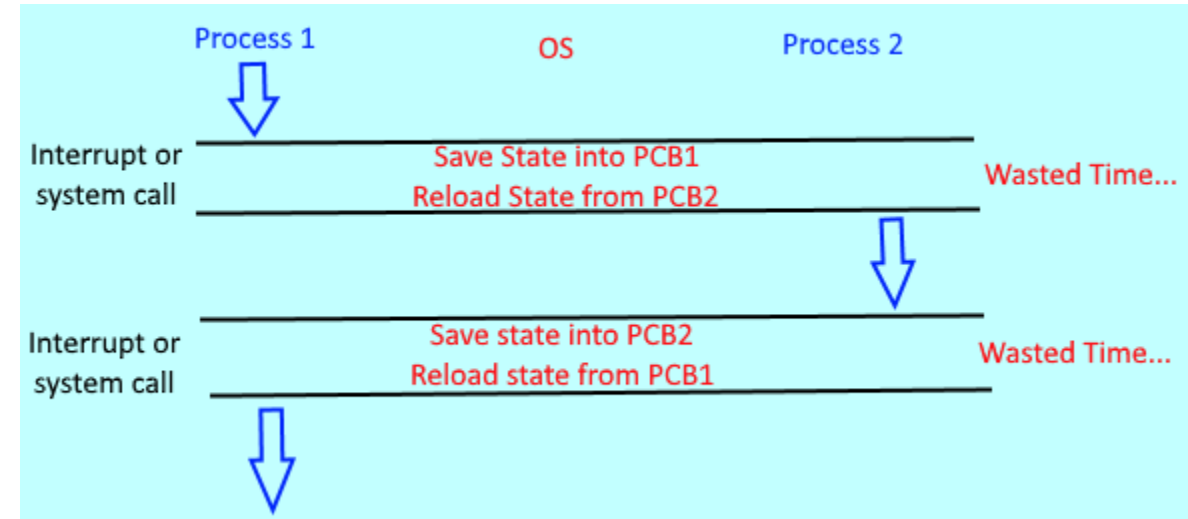
Process Context Switch

❑ Overhead:

- The time incurred by a context switch. The CPU does no useful work while switching.
- Translation lookaside buffer (TLB) flushing.

❑ Tradeoff:

- Context switch can improve the CPU utilization; however, context switch may incur overheads, which may reduce CPU utilization.



Process Context Switch

- ❑ When to do context switch
 - I/O system call
 - Time slice is up (timer interrupt)
 - Arrival of higher priority process
 - Yield by process itself (yield())
 - Termination of process
- ❑ Difference between context switch and mode switch?
 - The CPU is executing the same process when mode switch occurs.
 - The CPU is executing a different process when context switch occurs.
- ❑ Not all the system calls trigger context switch.