

ECE437/CS481

MO2B: PROCESSES & THREADS

PROCESSES CREATION

Chapter 3.3

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a thin line, dipping into a V-shape, and then rising back to a thin line, creating a stylized horizon or wave effect.

Process Creation & Termination

❑ New processes are created by the existing processes

- Creating process is called **parent**
- Created process is called **child**
- Linux: do **ps**, look for PID field

```
home@VirtualBox:~$ ps ux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
home     1114  0.0  0.8  46548  8512 ?        Ssl   Sep03   0:00 gnome-sess
home     1151  0.0  0.0   3856   140 ?        Ss    Sep03   0:00 /usr/bin/s
home     1154  0.0  0.0   3748   484 ?        S     Sep03   0:00 /usr/bin/d
home     1155  0.1  0.2   6656  3036 ?        Ss    Sep03   0:18 //bin/dbus
home     1157  0.0  0.2   9148  2368 ?        S     Sep03   0:00 /usr/lib/g
home     1162  0.0  0.2  31588  2296 ?        Ssl   Sep03   0:00 /usr/lib/g
home     1174  0.0  1.4 132472 14884 ?        Sl    Sep03   0:03 /usr/lib/g
```

❑ Each process can create/spawn child processes, and each process has parent except for the **top-most process**.

❑ Who and when creates **top-most process**?

- **init/systemd** is the **top-most process** with PID=1. It's created directly by the **kernel**.
- All other processes are either started directly by **init/systemd** or by one of its child processes.

Process Creation & Termination

- ❑ Spawn a new process via `fork()` system call
 - A child process get a copy of the address space of its parent process.
 - Both processes continue execution at the instruction after the `fork()`.
 - If `fork()` succeeds:
 - ✓ return 0 to the child
 - ✓ return child's PID to the parent
 - If `fork()` fails:
 - ✓ no child process is created, and it returns -1 in the parent process

} The `fork()` system
call "returns twice"

❑ Synopsis/Syntax of fork()

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork(void);
```

❑ Template for using fork()

```
pid_t pid;
pid=fork();
if (pid== 0) {
/* put code for child here */ }
else if (pid < 0) {
/* fork failed, put error handling here */ }
else {
/* fork successful; put remaining code for parent here */ }
```

Process Creation & Termination

❑ Example of fork()

```
#include <unistd.h>
#include <sys/types.h>
#include <stdio.h>

int main()
{
    pid_t pid = fork();
    if (pid == 0)
        printf("\tI'm the child process, and I received %d\n", pid);

    else
        printf("I'm the parent process, and my child PID is %d\n", pid);
    return 0;
}
```

Output:

I'm the parent process, and my child PID is 2172
I'm the child process, and I received 0

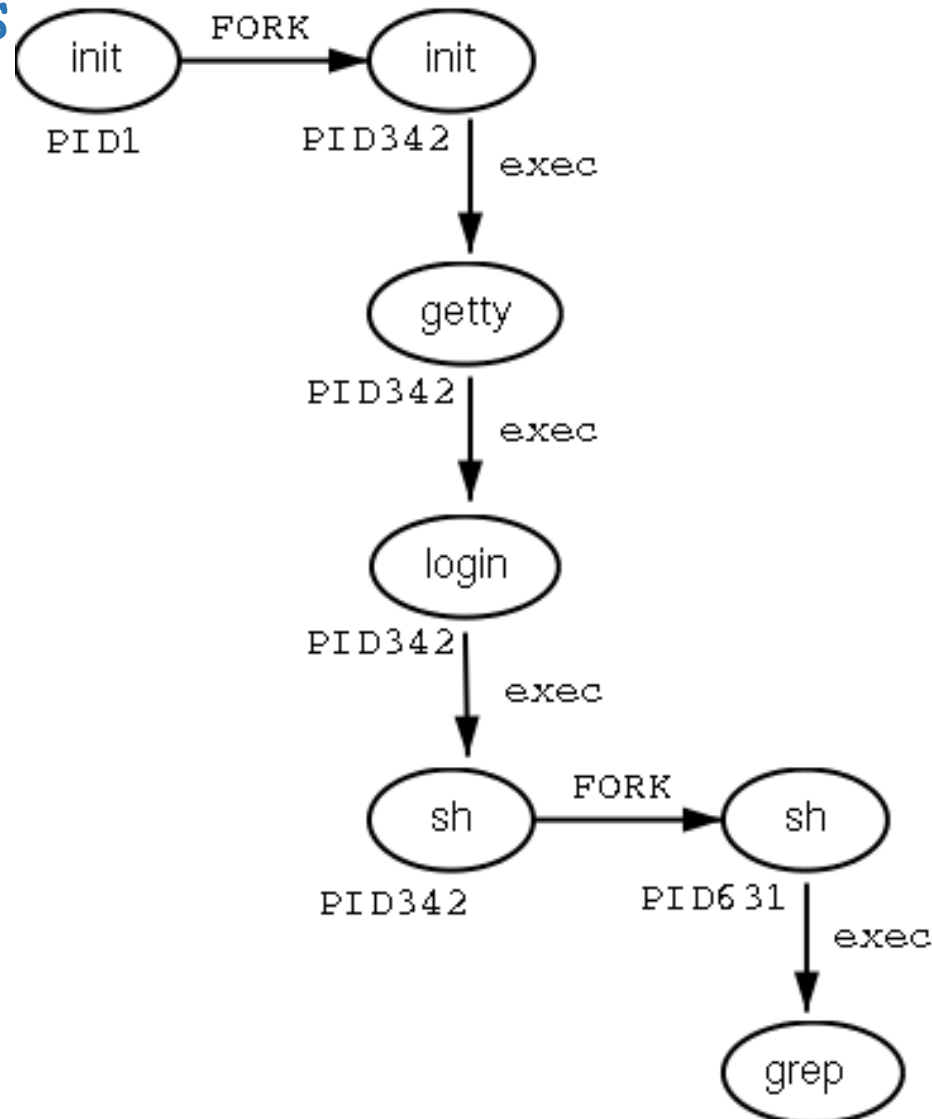
Process Creation & Termination

- ❑ Spawn a new process via `execve()` system call
 - `execve()` is called to let the new process **overwrite the calling process totally**
 - The address space of the calling process is reinitialized
 - There is **no return from a successful call to `execve()`**
 - ✓ return -1, if `execve()` fails
 - An `execve()` call often follows `fork()` to create a new process that runs another program/function.
 - ✓ Process A (the parent process) calls `fork()` to create a child Process B;
 - ✓ Process B immediately calls `execve()` to run a new program/function.

Process Creation & Termination

□ Example of process creations

- ✓ `fork()` creates a child process with a different PID.
- ✓ `execve()` overwrites a process with the same PID.



□ Synopsis/Syntax of `execve()`

```
#include <unistd.h>
```

```
int execve(const char *pathname, char *const argv[], ..., char *const envp[]);
```

- ✓ use "man -s 2 execve" to see more details
- ✓ Both `argv` and `envp` must end with a null pointer.

Process Creation & Termination

□ Example of `execve()`

```
#include <stdlib.h>
#include <unistd.h>
#include <stdio.h>

int
main(int argc, char **argv) {
    char *args[] = {"/bin/ls", "-m", 0}; /* each element represents a command line argument */
    char *env[] = { 0 }; /* leave the environment list null */
    printf("About to run /bin/ls\n");
    execve("/bin/ls", args, env);
    perror("execve"); /* if we get here, execve failed */
    exit(1);
}
```

Output:

```
shaun@shaun-VirtualBox:~$ ./test1
About to run /bin/ls
Desktop, Documents, Downloads, Music, Pictures, Public, Templates, Videos,
a.out, examples.desktop, hello, hello.c, test.c, test1, test1.c
```

Process Creation & Termination

❑ Terminate a process via `kill()` system call

- a process may terminate other process if it has such a privilege
- terminate another process by `kill()` system call, which sends a signal specified by PID.

❑ Synopsis/Syntax of `kill()`

```
#include <unistd.h>
#include <signal.h>
int kill(pid_t pid, int sig);
```

❑ Use "man -s 2 kill" to see more details

Number	Name (short name)	Description	Used for
0	SIGNULL (NULL)	Null	Check access to pid
1	SIGHUP (HUP)	Hangup	Terminate; can be trapped
2	SIGINT (INT)	Interrupt	Terminate; can be trapped
3	SIGQUIT (QUIT)	Quit	Terminate with core dump; can be trapped
9	SIGKILL (KILL)	Kill	Forced termination; cannot be trapped
15	SIGTERM (TERM)	Terminate	Terminate; can be trapped
24	SIGSTOP (STOP)	Stop	Pause the process; cannot be trapped. This is default if signal not provided to kill command.
25	SIGTSTP (STP)	Terminal	Stop/pause the process; can be trapped
26	SIGCONT (CONT)	Continue	Run a stopped process

Process Creation & Termination

❑ Example of kill()

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>

int main(void){
    pid_t retVal;
    retVal = fork();
    if(retVal > 0){
        int i = 0;
        while(i++ < 5){
            printf("in the parent process.\n");
            sleep(1);
        }
        //kill the child process
        kill(retVal, SIGKILL);
    } else if (retVal == 0){
        int i = 0;
        //will not ever get to 15, because
        //the parent process will kill it
        while(i++ < 15){
            printf("In the child process.\n");
            sleep(1);
        }
    } else {
        //something bad happened.
        printf("Something bad happened.");
        exit(EXIT_FAILURE);
    }
    return 0;
}
```


Output:

```
shaun@shaun-VirtualBox:~$ ./test_wait
in the parent process.
In the child process.
in the parent process.
In the child process.
in the parent process.
In the child process.
in the parent process.
In the child process.
in the parent process.
In the child process.
```

Process Creation & Termination

❑ Wait for termination of a process via `wait()` system call

- A parent process can temporarily suspend its execution in order to let its child process to be execute.
- After child process **terminates**, parent continues its execution after waiting system call instruction.
- Child process may terminate due to any of these:
 - ✓ It calls `exit()`;
 - ✓ It returns (an int) in `main()`;
 - ✓ It receives a signal (from the OS or another process) whose action is to terminate.



Process Creation & Termination

- ❑ Wait for termination of a process via `wait()` system call
 - If **only one child process was created**, then `wait()` suspends the parent and returns process ID of the child process once it is terminated.
 - If **more than one child processes were created**, then `wait()` suspends the parent process until one of its child processes terminates.
 - If the process has no child process, then `wait()` returns -1.
- ❑ Synopsis/Syntax of `wait()`

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

Process Creation & Termination

❑ Example of wait()

```
#include<stdio.h>
#include<sys/wait.h>
#include<unistd.h>

int main()
{
    pid_t pid=fork();
    if (fork()== 0)
        printf("Child is running\n");
    else
    {
        printf("Parant is running\n");
        wait(NULL);
        printf("Child has been terminated\n");
    }

    return 0;
}
```



```
#include<stdio.h>
#include<sys/wait.h>
#include<unistd.h>

int main()
{
    pid_t pid1=fork();
    pid_t pid2=fork();
    if (pid2==0)
    {
        printf("Child is running\n");
    }
    else
    {
        printf("Parent is running\n");
        wait(NULL);
        printf("Child has been terminated\n");
    }
    return 0;
}
```

Process Creation & Termination

❑ Example of wait()

```
#include<stdio.h>
#include<sys/wait.h>
#include<unistd.h>

int main()
{
    pid_t pid1=fork();
    pid_t pid2=fork();
    if (pid2==0)
    {
        printf("Child is running\n");
    }
    else
    {
        printf("Parent is running\n");
        wait(NULL);
        printf("Child has been terminated\n");
    }
    return 0;
}
```

Output:

```
shaun@shaun-VirtualBox:~/OS_code/proc$ ./test_wait1
Parent is running
Child is running
Child has been terminated
shaun@shaun-VirtualBox:~/OS_code/proc$ Parent is running
Child is running
Child has been terminated
█
```

Process Creation & Termination

❑ Child status information in wait(*status)

- Status indicates the **termination info** of the child process reported by wait.
- **Termination info**: **encode** both the termination status of the process as well as the reason of termination.
- In order to find the termination info, **WIF... macros** is applied. For example,
 1. WIFEXITED(status): returns true if child is terminated by exit().
WEXITSTATUS(status): return the parameter of exit() if WIFEXITED(status) is true.
 2. WIFSIGNALED(status): returns true if the child process was terminated by a signal.
WTERMSIG(status): gives the index of the terminating signal if WIFSIGNALED(status) is true.
- Find more info about **WIF... macros** via <http://man7.org/linux/man-pages/man2/waitpid.2.html>

Process Creation & Termination

❑ Example of wait(&status)

```
#include<stdio.h>
#include<stdlib.h>
#include<sys/wait.h>
#include<unistd.h>

void waitexample()
{
    int stat;
    if (fork() == 0)
        exit(1);
    else
        wait(&stat);
    if (WIFEXITED(stat))
        printf("Exit status: %d\n", WEXITSTATUS(stat));
    else if (WIFSIGNALED(stat))
        psignal(WTERMSIG(stat), "Exit signal");
}

int main()
{
    waitexample();
    return 0;
}
```

Output:

```
shaun@shaun-VirtualBox:~$ ./test_wait2
Exit status: 1
```

Process Creation & Termination

❑ Wait for termination of a specific process via **waitpid()** system call

❑ Synopsis/Syntax of waitpid()

```
#include <sys/wait.h>
```

```
pid_t waitpid(child_pid, *status, options);
```

➤ **options:**

- 1) If 0, it means no option, i.e., the parent has to wait for the termination of the target child process. Return: target child process PID.
- 2) If WNOHANG, it means parent does not wait if the target child process does not terminate. (not block parent process). Return: target child process PID if it is terminated, or 0 if it is not terminated.
- 3) If WUNTRACED, waitpid() **also** returns if the target child process has stopped.
- 4) If WCONTINUED, waitpid() **also** returns if the stopped child process has been resumed by delivery of SIGCONT.

➤ Find more info about waitpid() via <http://man7.org/linux/manpages/man2/waitpid.2.html>

Process Creation & Termination

❑ Example of waitpid()

```
#include<stdio.h>
#include<stdlib.h>
#include<sys/wait.h>
#include<unistd.h>

void waitexample()
{
    int i, stat;
    pid_t pid[5];
    for (i=0; i<5; i++)
    {
        if ((pid[i] = fork()) == 0)
        {
            sleep(1);
            exit(100 + i);
        }
    }

    // Using waitpid() and printing exit status
    // of children.
    for (i=0; i<5; i++)
    {
        pid_t cpid = waitpid(pid[i], &stat, 0);
        if (WIFEXITED(stat))
            printf("Child %d terminated with status: %d\n",
                   cpid, WEXITSTATUS(stat));
    }
}

int main()
{
    waitexample();
    return 0;
}
```

Output:

```
shaun@shaun-VirtualBox:~$ ./test_waitpid
child 2701 terminated with status: 100
child 2702 terminated with status: 101
child 2703 terminated with status: 102
child 2704 terminated with status: 103
child 2705 terminated with status: 104
```


© by Dr. X. Sun

Process Creation & Termination

❑ Zombie process

- When a process terminates, its resources are deallocated by the operating system. However, its entry in the **process table** as well as its PCB are still maintained until its parent calls `wait()`.
- A **zombie process** is a process that has terminated, but its parent has not yet called `wait()`.
- Zombie processes only occupy few system resources (to store their PCBs). However, each zombie process retains its process ID (PID), and total number of PIDs is limited.
 - ✓ Consuming all the available PID pool, thus preventing creating new processes.

```
int main()
{
    // Fork returns process id
    // in parent process
    pid_t child_pid = fork();

    // Parent process
    if (child_pid > 0)
        sleep(50);

    // Child process
    else
        exit(0);

    return 0;
}
```

Process Creation & Termination

❑ Zombie process

- In the Ubuntu OS system, if a process that is still existed in the system, but its parent has terminated and has not called wait(). This process is also considered a **zombie process**.

```
#include <stdio.h>
#include <unistd.h>

int main ()
{
    pid_t child_pid;
    child_pid = fork();
    if (child_pid > 0) {
        printf ("child process ID %d\n",child_pid);
        exit(0); }
    else {
        printf("I am the child process before sleep.\n");
        sleep(20);
        printf("I am the child process after sleep.\n");
        exit (0);}
}
```


Process Creation & Termination

❑ Zombie process

- In the Ubuntu OS system, if a process that is still existed in the system, but its parent has terminated and has not called wait(). This process is also considered a **zombie process**.

```
shaun 5971 0.0 0.2 581628 2428 ? Ssl 10:05 0:00 /usr/libexec/gnome-shell-calendar-server
shaun 5977 0.0 0.2 464156 2448 ? Ssl 10:05 0:00 /usr/libexec/evolution-source-registry
shaun 5985 0.0 0.3 707748 3652 ? Ssl 10:05 0:00 /usr/libexec/evolution-calendar-factory
shaun 5995 0.0 0.2 156200 2840 ? Sl 10:05 0:00 /usr/libexec/dconf-service
shaun 6000 0.0 0.3 681664 3692 ? Ssl 10:05 0:00 /usr/libexec/evolution-addressbook-factory
shaun 6018 0.0 0.3 317548 3304 ? Sl 10:05 0:00 /usr/libexec/gvfsd-trash --spawner :1.3 /org/gtk/gvfs/exec_sp
shaun 6023 0.0 0.1 2598772 1912 ? Sl 10:05 0:00 /usr/bin/gjs /usr/share/gnome-shell/org.gnome.Shell.Notificat
shaun 6036 0.0 0.2 313920 2584 ? Ssl 10:05 0:00 /usr/libexec/gsd-a11y-settings
shaun 6038 0.0 0.6 570392 6936 ? Ssl 10:05 0:00 /usr/libexec/gsd-color
shaun 6040 0.0 0.2 374364 2588 ? Ssl 10:05 0:00 /usr/libexec/gsd-datetime
shaun 6041 0.0 0.2 314164 2968 ? Ssl 10:05 0:00 /usr/libexec/gsd-housekeeping
shaun 6042 0.0 0.5 348352 5300 ? Ssl 10:05 0:00 /usr/libexec/gsd-keyboard
shaun 6046 0.0 0.9 690548 9096 ? Ssl 10:05 0:00 /usr/libexec/gsd-media-keys
shaun 6047 0.0 0.6 348772 6140 ? Ssl 10:05 0:00 /usr/libexec/gsd-power
shaun 6051 0.0 0.3 248100 2996 ? Ssl 10:05 0:00 /usr/libexec/gsd-print-notifications
shaun 6053 0.0 0.3 457100 3060 ? Ssl 10:05 0:00 /usr/libexec/gsd-rfkill
shaun 6054 0.0 0.2 235508 2468 ? Ssl 10:05 0:00 /usr/libexec/gsd-screensaver-proxy
shaun 6057 0.0 0.3 469040 3640 ? Ssl 10:05 0:00 /usr/libexec/gsd-sharing
shaun 6060 0.0 0.2 318080 2536 ? Ssl 10:05 0:00 /usr/libexec/gsd-smartcard
shaun 6063 0.0 0.2 321784 2472 ? Ssl 10:05 0:00 /usr/libexec/gsd-sound
shaun 6065 0.0 0.2 387884 2464 ? Ssl 10:05 0:00 /usr/libexec/gsd-usb-protection
shaun 6073 0.0 0.4 274188 4908 ? Ssl 10:05 0:00 /usr/libexec/gsd-wacom
shaun 6075 0.0 0.2 318248 2308 ? Ssl 10:05 0:00 /usr/libexec/gsd-wwan
shaun 6079 0.0 0.6 348996 6156 ? Ssl 10:05 0:00 /usr/libexec/gsd-xsettings
shaun 6093 0.0 0.3 789816 3780 ? Sl 10:05 0:00 /usr/libexec/evolution-data-server/evolution-alarm-notify
shaun 6113 0.0 0.1 231800 1352 ? Sl 10:05 0:00 /usr/libexec/gsd-disk-utility-notify
shaun 6145 0.0 0.2 342180 2420 ? Sl 10:05 0:00 /usr/libexec/gsd-printer
shaun 6191 0.0 0.1 166548 1668 ? Sl 10:05 0:00 /usr/libexec/ibus-engine-simple
shaun 6254 0.0 0.3 162296 3456 ? Ssl 10:06 0:00 /usr/libexec/gvfsd-metadata
shaun 6257 0.0 1.4 496532 14404 ? Sl 10:06 0:00 update-notifier
shaun 6293 0.3 3.4 821124 34036 ? Rsl 10:06 0:07 /usr/libexec/gnome-terminal-server
shaun 6311 0.0 0.3 10612 3196 pts/0 Ss+ 10:06 0:00 bash
shaun 6456 0.2 3.2 991712 32448 ? Sl 10:07 0:05 /usr/bin/nautilus --gapplication-service
shaun 6498 0.0 0.3 317740 3312 ? Sl 10:09 0:00 /usr/libexec/gvfsd-network --spawner :1.3 /org/gtk/gvfs/exec_
shaun 6504 0.0 0.6 450336 6048 ? Sl 10:09 0:00 /usr/libexec/gvfsd-smb-browse --spawner :1.3 /org/gtk/gvfs/ex
root 7287 0.0 0.0 0 0 ? I< 10:10 0:00 [xfsalloc]
root 7288 0.0 0.0 0 0 ? I< 10:10 0:00 [xfs_mru_cache]
root 7292 0.0 0.0 0 0 ? S 10:10 0:00 [jfsIO]
root 7293 0.0 0.0 0 0 ? S 10:10 0:00 [jfsCommit]
root 7294 0.0 0.0 0 0 ? S 10:10 0:00 [jfsSync]
shaun 8060 0.2 4.9 820392 49232 ? Sl 10:11 0:03 /usr/bin/gedit --gapplication-service
shaun 8362 0.0 0.4 10612 4756 pts/1 Ss 10:13 0:00 bash
root 8435 0.0 0.0 0 0 ? T 10:30 0:00 [kworker/u2:1-events_unbound]
shaun 8465 0.0 0.0 2496 72 pts/0 S 10:36 0:00 ./Zombie
shaun 8466 0.0 0.3 11496 3172 pts/1 R+ 10:36 0:00 ps aux
shaun@shaun-VirtualBox:~$
```

Zombie process
with PID=8465

```
Zombie.c:16:2: warning: incompatible implicit declaration of built-in function
'exit'
16 |     exit (0);
    |     ^~~~~
Zombie.c:16:2: note: include '<stdlib.h>' or provide a declaration of 'exit'
shaun@shaun-VirtualBox:~/OS$ ./Zombie
child process ID 8465
shaun@shaun-VirtualBox:~/OS$ I am the child process before sleep.
I am the child process after sleep.
^
```

Process Creation & Termination

❑ How to kill a zombie process

- **Solution_1:** a child process will automatically send a **SIGCHLD** signal to the parent process after its termination. Once receiving the signal, the parent process could execute the `wait()` system call and clean up its zombie children.

Assumption: the parent process should be programmed properly (i.e., invoke `wait()`) when **SIGCHLD** signals are trapped.

- **Solution_2:** kill the zombie's parent process.
 - ✓ All its child process (including zombie process) become **orphan processes**.
 - ✓ Init process (`pid=1`) becomes the new parent to **orphan processes**.
 - ✓ Init process periodically executes the `wait()` system call to clean zombie process.

Process Creation & Termination

□ More about Unix/Linux Process's Attributes

➤ PID

- ✓ each process has a non-negative, system-unique # as its PID
- ✓ some fixed process IDs, such as PID=1, init process
- ✓ obtain its own process ID by calling `getpid(void)`
- ✓ obtain its parent process ID by calling `getppid(void)`

Process Creation & Termination

□ More about Unix/Linux Process's Attributes

➤ Process group ID (GID)

- ✓ process group == a set of related processes
process group is used to control the distribution of signals.
- ✓ when the process is created, it becomes a member of the process group of its parent.
- ✓ process group ID (GID)=PID of the process group leader (normally, the first member of the process group)
- ✓ a process finds another process *p*'s GID using `getpgid(p)`
- ✓ put a process into a process group using `setpgid(pid, pgid)`
- ✓ create a new process group with process group leader pid using `setpgid(pid, 0)`
- ✓ `setpgrp()=setpgid(0, 0)`, which is to create a new process with process group leader equaling to the calling process.

Process Creation & Termination

□ More about Unix/Linux Process's Attributes

➤ Session ID (SID)


- ✓ session = a set of process groups sharing a control terminal.
- ✓ when a process is created, it becomes a member of its parent's session.
- ✓ the session ID of a session = the process ID of the first member of the session.
- ✓ a process finds the ID of its session using the system call `getsid()`.
- ✓ start a new session by calling `setsid(void)` (return the new SID)

Process Creation & Termination

□ More about Unix/Linux Process's Attributes

➤ pipe() system call

- ✓ pipe is one-way communication (one process write to the pipe, and the other process reads from the pipe). If a pipe is created, there is an area of main memory that is treated as a "virtual file". One process can write to this "virtual file" and another related process can read from it later on.
- ✓ If a process attempts to read from an empty pipe, the process is block.



Process Creation & Termination

□ Example of process group and pipe()

```
#define _POSIX_SOURCE
#include <unistd.h>
#include <sys/types.h>
#include <stdio.h>

main() {
    pid_t pid;
    int p1[2], p2[2];
    char c='?';

    if (pipe(p1) != 0)
        perror("pipe() #1 error");
    else if (pipe(p2) != 0)
        perror("pipe() #2 error");
    else
        if ((pid = fork()) == 0) {
            printf("child's process group id is %d\n", (int) getpggrp());
            write(p2[1], &c, 1);
            read(p1[0], &c, 1);
            puts("child is waiting for parent to complete task");
            printf("child's process group id is now %d\n", (int) getpggrp());
            exit(0);
        }
        else {
            printf("parent's process group id is %d\n", (int) getpggrp());
            read(p2[0], &c, 1);
            printf("parent is performing setpgid() on pid %d\n", (int) pid);
            if (setpgid(pid, 0) != 0)
                perror("setpgid() error");
            write(p1[1], &c, 1);
            printf("parent's process group id is now %d\n", (int) getpggrp());
            sleep(5);
        }
}
```

```
shaun@shaun-VirtualBox:~$ ./test_setpgrp
parent's process group id is 2722
child's process group id is 2722
parent is performing setpgid() on pid 2723
parent's process group id is now 2722
child is waiting for parent to complete task
child's process group id is now 2723
```