

ECE437/CS481

M03A: CPU SCHEDULING
SCHEDULING CONCEPT & ALGORITHMS

CHAPTER 6.1-6.3

Xiang Sun

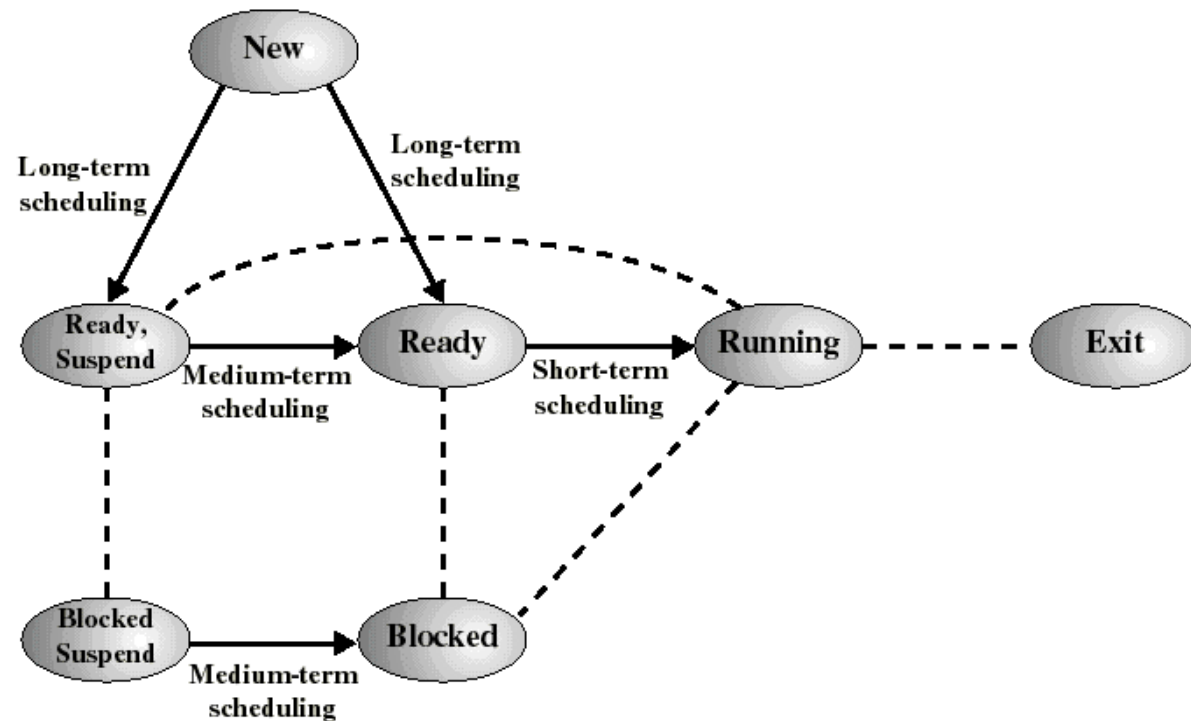
The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting from the left edge, dipping down in the center, and rising back up to the right edge, creating a stylized wave or valley shape.

Levels of Scheduling

□ Long-term scheduling/Job scheduler (High-level)

- A long-term scheduler determines which programs are admitted to the system for processing (i.e., loaded into the main memory).
- Coarse-grained control of the **degree of multiprogramming**
- Should balance different types of processes:
 - **I/O-intensive process**: spends more time doing I/O than computations, many short CPU bursts.
 - **CPU-intensive process**: spends more time doing computations; few very long CPU bursts.



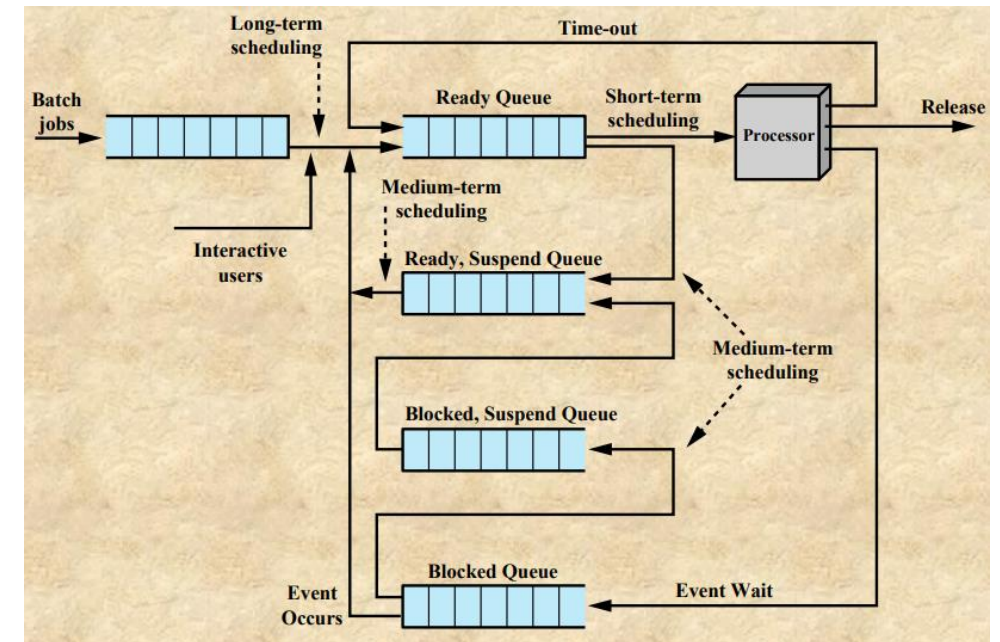
Levels of Scheduling

❑ Medium-term scheduling (swapping in/out)

- Adjust the degree of multiprogramming by **swapping**.
- Swapping: removes a process from memory, stores on disk, and bring back in from disk to continue execution later on.

- ✓ **Swap out**: memory-to-disk
- ✓ **Swap in**: disk-to-memory

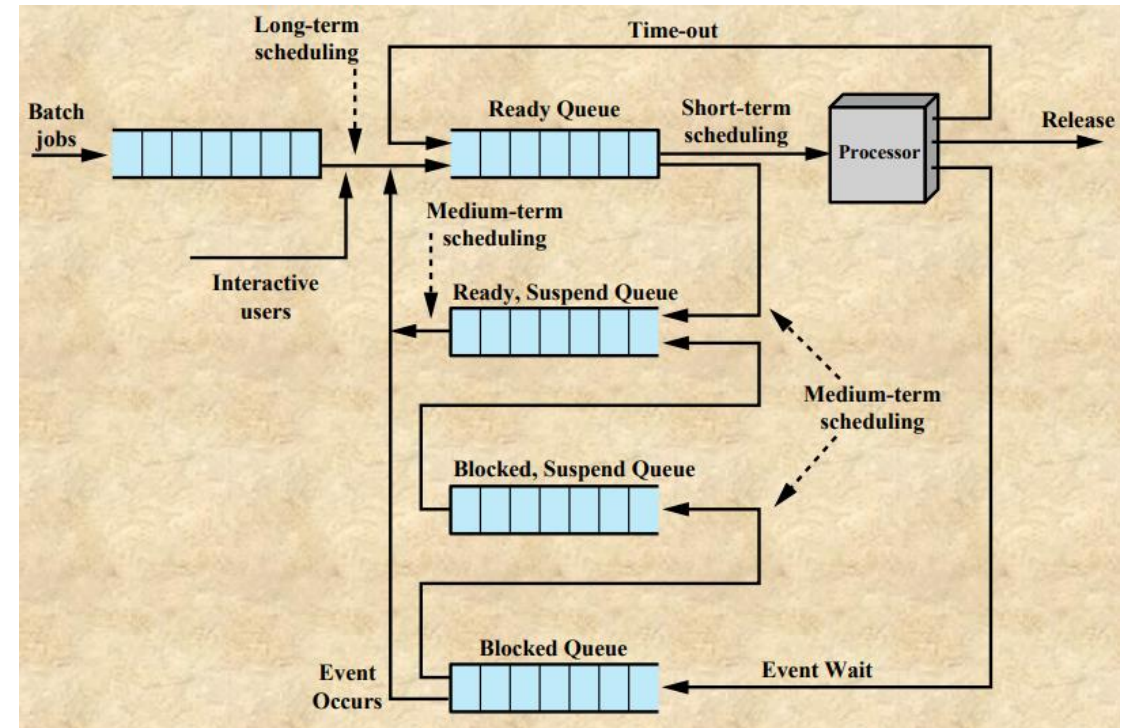
- Normally, the medium-term scheduler may decide to swap out a process which has not been active for some time, or a process which has a low priority, or a process which is taking up a large amount of memory in order to free up main memory for other processes, etc.



Levels of Scheduling

❑ Short-term scheduling/CPU scheduling

- Determine the execution order of the processes in the ready queue.
- Invoked when an event occurs that may lead to the blocking/termination of the current process or that may provide an opportunity to preempt a currently running process in favor of another.
 - ✓ Clock interrupts
 - ✓ I/O interrupts
 - ✓ System calls
 - ✓ Signals, etc.
- **Scheduler is consuming CPU too!** It executes most frequently, must be very careful about its computation overhead.



Scheduling Objectives

□ Metrics—system-wide

- Maximize processor/CPU utilization
 - ✓ percentage of time that CPU is running users' processes, to keep system as busy as possible
- Maximize throughput
 - ✓ number of processes completed per time unit
 - ✓ number of instructions executed per time unit
- Fairness
 - ✓ don't starve any processes—treat the all the same



Scheduling Objectives

- ❑ Metrics—per process/user-oriented
 - Minimize waiting time in the **ready queue**
 - Minimize real time/wall clock time
 - ✓ equal to sum of the waiting time in the ready and waiting queue, plus running time.
 - Minimize response time
 - ✓ for interactive job, time from the submission of a request until **the first response happens**.



Scheduling Objectives

□ Metrics

- Achieve a balance between response time and utilization
- Minimize overhead (system level)
 - ✓ Context switching
 - ✓ Scheduling complexity
- It is difficult to find the optimal solution of the scheduling since the scheduling problem is mostly an **NP-hard/NP-complete** problem (e.g., job shop scheduling). Instead, looking for heuristic approaches is the common way.



❑ Non-preemption V.S. Preemption

➤ Non-preemption

- ✓ The running process continues execution until either it terminates/blocked/yield, even though a new process is scheduled by the scheduling algorithm.

➤ Preemption

- ✓ The running process is immediately suspended when a new process is scheduled by the scheduling algorithm.

❑ Concurrency V.S. Parallelism

➤ Concurrency

- ✓ schedule multiple processes onto a single CPU—time multiplex manner

➤ Parallelism

- schedule multiple processes onto multiple CPUs—spatial multiplex

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch

❑ Decision mode

- non-preemption
- preemption—high **priority** or periodically (the process's time slice expires)

❑ Priority function

- static information, e.g., CPU-intensive or I/O intensive, memory requirement, service time
- dynamic information, e.g., relative deadline, recent CPU consumption

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch

□ Events affect/trigger the scheduling algorithm:

- 1) Current process goes from **running to waiting state** (e.g., wait for I/O)
- 2) Current process terminates, **running to termination state**
- 3) Current process goes from **running to ready state** (e.g., time slice is up or yield)

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch


- ❑ Dispatcher module gives control of the CPU to the process selected by the **short-term scheduler**; this involves:
 - 1) switching context
 - 2) switching between user mode and kernel mode
- ❑ **Dispatch latency**: time consumption of the dispatcher to stop one process and start another.

Scheduling Algorithms

❑ First-Come, First-Served (FCFS)

- General specification
 - ✓ Decision mode: non-preemption
 - ✓ Priority: arrival time
- As a process become ready, it join the ready queue, **scheduler always selects process from the front of the ready queue.**

Process	Arrival time	Service time (Burst time)
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Waiting time for P1 = 0; P2 = 1; P3 = 5, P4 = 7, P5 = 10,
Average waiting time: $(0 + 1 + 5 + 7 + 10)/5 = 4.6$

Scheduling Algorithms

❑ First-Come, First-Served (FCFS)

- Simple to implement; low overhead, since no priority calculation, no extra context switch.
- Average waiting time may be long and suffer from **convoy effect**.
 - ✓ Convoy effect: long waiting time due to the slow processes.

Case 1:

- ✓ Suppose that there are three processes already in the ready queue, and their arrival order: P₁, P₂, P₃.

✓ The Gantt Chart for the schedule is



- ✓ Waiting time P₁=0; P₂=24; P₃=27
- ✓ Average waiting time: $(0+24+27)/3=17$

Process	Service time
P ₁	24
P ₂	3
P ₃	3

Case 2:

- ✓ Suppose that there are three processes already in the ready queue, and their arrival order: P₂, P₃, P₁.

✓ The Gantt Chart for the schedule is



- ✓ Waiting time P₁=6; P₂=0; P₃=3
- ✓ Average waiting time: $(6+0+3)/3=3$

Scheduling Algorithms

❑ Shortest Job Next (SJN)/Shortest Job First(SJF)

➤ General specification

- ✓ Decision mode: **non-preemption by default** (could be implemented as preemption)
- ✓ Priority: $1/\text{service time}$ —a process with a shorter service time has a higher priority

➤ Pros:

- ✓ Decrease the average waiting time as compared to FCFS

➤ Cons:

- ✓ Higher complexity as compared to FCFS
- ✓ Difficult to **predict service time**
- ✓ Risk to **starve** slow process, as long as there are fast processes around

❑ Shortest Remaining-time First (SRF)

➤ General specification


- ✓ Decision mode: **preemption**
- ✓ Priority: $1/(\text{service time} - \text{running time})$

- Still needs to estimate service time and can potentially starve slow processes.
- May give less average waiting time than SJN/SJF but lead to more context switch.

Scheduling Algorithms

❑ Shortest Job Next (SJN)/Shortest Job First(SJF)

- Suppose that there are four processes already in the ready queue.
- The Gantt Chart for the schedule is:




Process	Service time
P1	6
P2	8
P3	7
P4	3

- Waiting time for P₁=3 , P₂=16 , P₃=9 , P₄=0
- Average waiting time: $(3+16+9+0)/4=7$
- Suppose FCFS is applied, average waiting time: $(0+6+14+21)/4=10.25$

Scheduling Algorithms

Shortest Job Next (SJN)/Shortest Job First(SJF)



Scheduling diagram

Process	Arrival time	Service time
P1	0	3
P2	2	6
P3	4	4
P4	6	5
P5	8	2

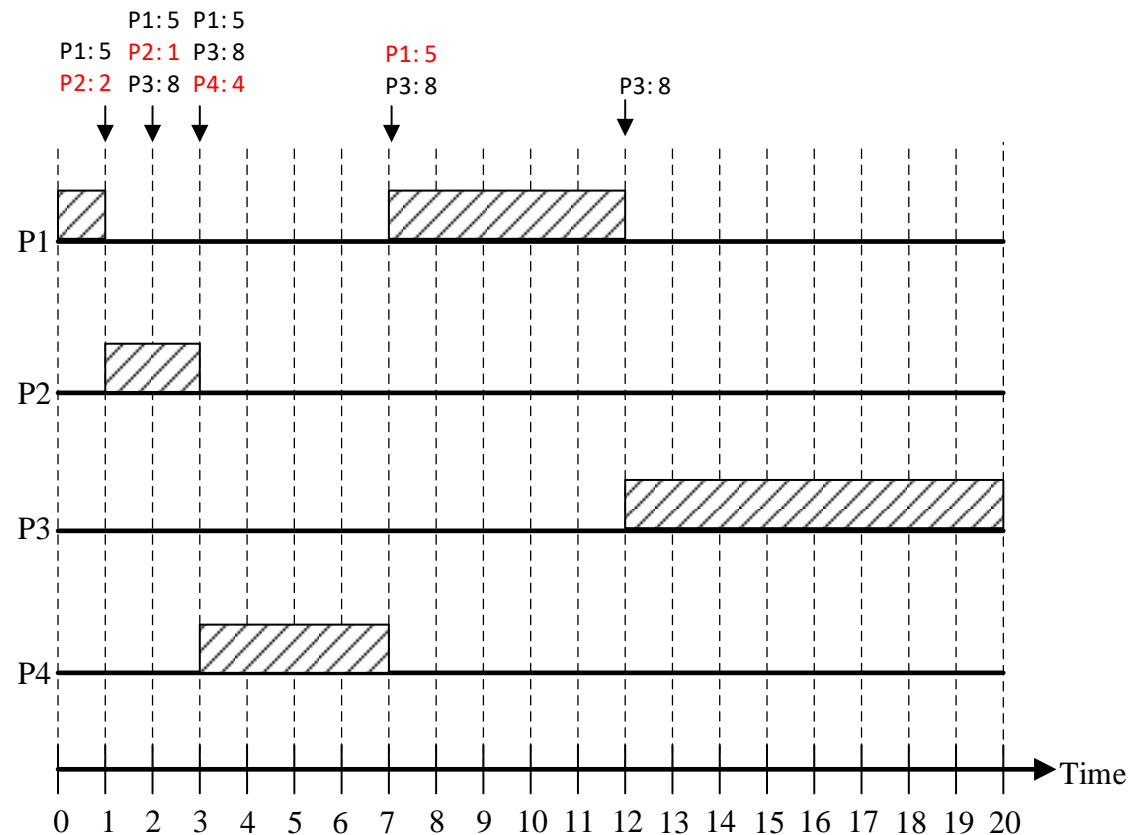
- Waiting time for P1=0 , P2=1, P3=7 , P4=9, P5=1.
- Average waiting time: $(0+1+7+9+1)/5=3.6$
- Suppose FCFS is applied, average waiting time: $(0+1+5+7+10)/5=4.6$

Scheduling Algorithms

❑ Shortest Remaining-time First (SRF) VS Shortest Job Next (SJN)

➤ Suppose we have four processes arriving based on the order showing in the table.

➤ The scheduling diagram for SRF is



Process	Arrival time	Service time
P1	0	6
P2	1	2
P3	2	8
P4	3	4

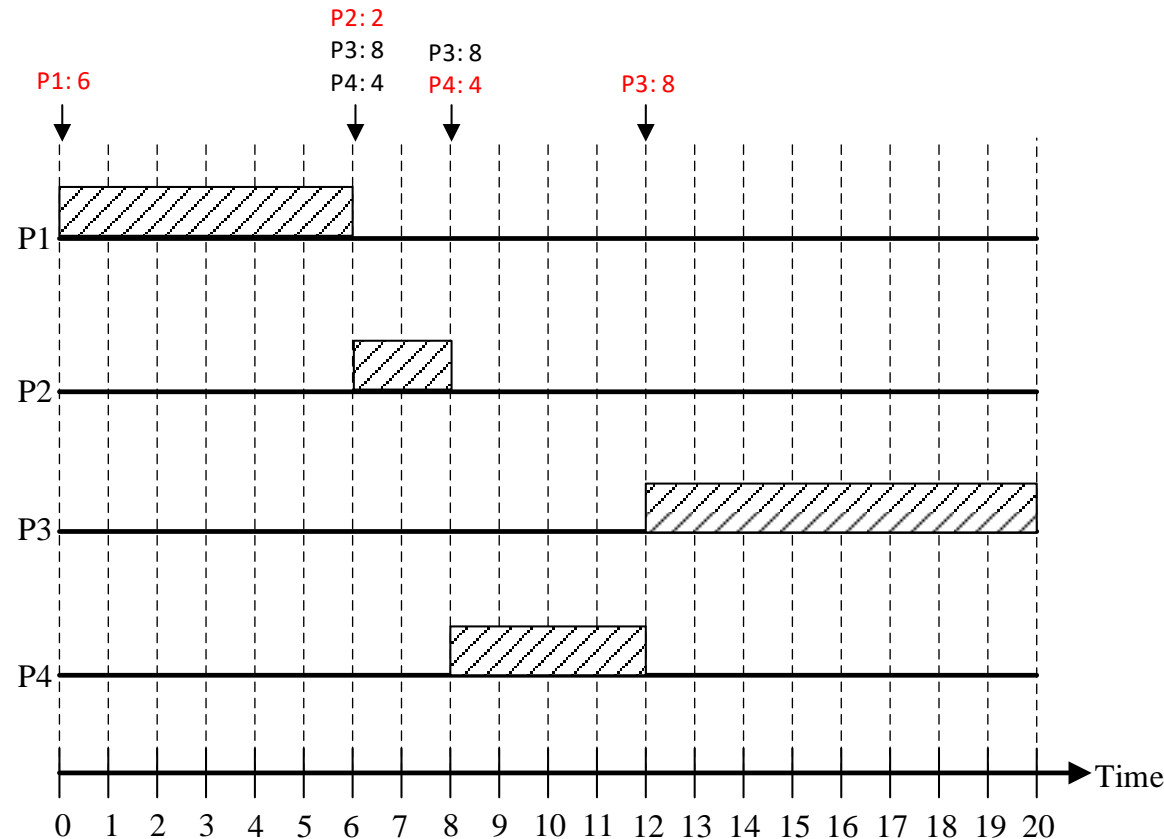
- The priorities of the processes are recalculated for each arrival.
- Waiting time for P1=6 , P2=0 , P3=10 , P4=0.
- Average waiting time: $(6+0+10+0)/4=4$

Scheduling Algorithms

❑ Shortest Remaining-time First (SRF) VS Shortest Job Next (SJN)

➤ Suppose we have four processes arriving in the order showing in the table.

➤ The scheduling diagram for SJN is



Process	Arrival time	Service time
P1	0	6
P2	1	2
P3	2	8
P4	3	4

- The priorities of the processes is static.
- The average waiting time: $(0+5+10+5)/4=5$

Scheduling Algorithms

□ Round Robin (RR)


- General specification
 - ✓ decision: preemption, periodically with time slice
 - ✓ priority: equal
- Use preemption based on clock - time slicing (time quantum), generate interrupt at periodic intervals.
- When an interrupt occurs, the running process is placed back to the end of Ready queue, the next process is selected to run based on FCFS.
- Designed especially for interactive jobs.
- What's the right length of a time slice?
 - ✓ short time slice leads to processes move through quickly, thus having high overhead to deal with process scheduling and context switching.
 - ✓ long time slice makes RR degenerate into FCFS.
 - ✓ should be slightly greater than average service time.

Scheduling Algorithms

□ Round Robin (RR)

- Suppose that there are three processes in the ready queue, and their arrival order: P1, P2, P3. The time quantum length is 4.

- The Gantt Chart for the schedule is:




Process	Service time
P1	24
P2	3
P3	3

- Wait time for P1=6; P2=4; P3=7
- Average waiting time: $(6+4+7)/3=5.67$.
 - ✓ Recall that the average waiting time for FCFS is 17.
 - ✓ However, average waiting time of RR < average waiting time of FCFS is not always true.
- The number of context switching in RR \geq The number of context switching in FCFS

Scheduling Algorithms

Round Robin (RR)- time quantum=1



Process	Arrival time	Service time (Burst time)
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

Waiting time for P1 = 1; P2 = 10; P3 = 9, P4= 9, P5 = 5,

Average waiting time: $(1 + 10 + 9 + 9 + 5)/5 = 6.8$

Recall that the average waiting time for applying FCFS is 4.6 in Slide 12, but waiting_time (RR) > waiting_time (FCFS) is not always true.