

ECE437/CS481

M03B: CPU SCHEDULING  
MORE ABOUT SCHEDULING

CHAPTER 6.4-6.8

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a small upward curve on the left, dipping into a V-shape in the center, and then curving back up on the right before continuing as a straight line to the edge.

# Dynamics with feedback

## □ Feedback in CPU scheduling

- Remaining time of a process may vary over time.
- Processes may be terminated and created by the system over time.
- Feedback is used to dynamic adjust the priorities of different processes.



# Dynamics with feedback

## □ Feedback to emphasize on **service time fairness**

- Dynamically bookkeeping two variables:
  - ✓ Required/Promised CPU time:  $T_p$
  - ✓ Actual CPU time used:  $T_u$
- Let priority function  $F = T_p/T_u$ 
  - ✓ The larger the value, the higher in priority
  - ✓ case  $F == 1$ , just right, kept promise.
  - ✓ case  $F > 1$ , under provisioned, the process can have more CPU time.
  - ✓ case  $F < 1$ , over provisioned, the process has to slow down.

# Dynamics with feedback

## □ Feedback to emphasize on **aging**

- Favor a process who spends more waiting time at the ready/waiting queue.
- Dynamically bookkeeping two variables:
  - ✓ Total waiting time:  $T_w$
  - ✓ Total time used:  $T_u$
- Let priority function  $F = T_w/T_u$ 
  - ✓ The larger the value, the higher in priority



# Classification of processes

## □ By the nature of requirement

### ➤ Foreground process

- ✓ Users can directly interact with a foreground process, for example, via a terminal/shell.
- ✓ Executing a foreground process does not disable the user to execute other background processes via the same terminal until it terminates.

### ➤ Background process

- ✓ Users cannot directly interact with a background process.
- ✓ Executing a background process does not disable the execution of other processes.

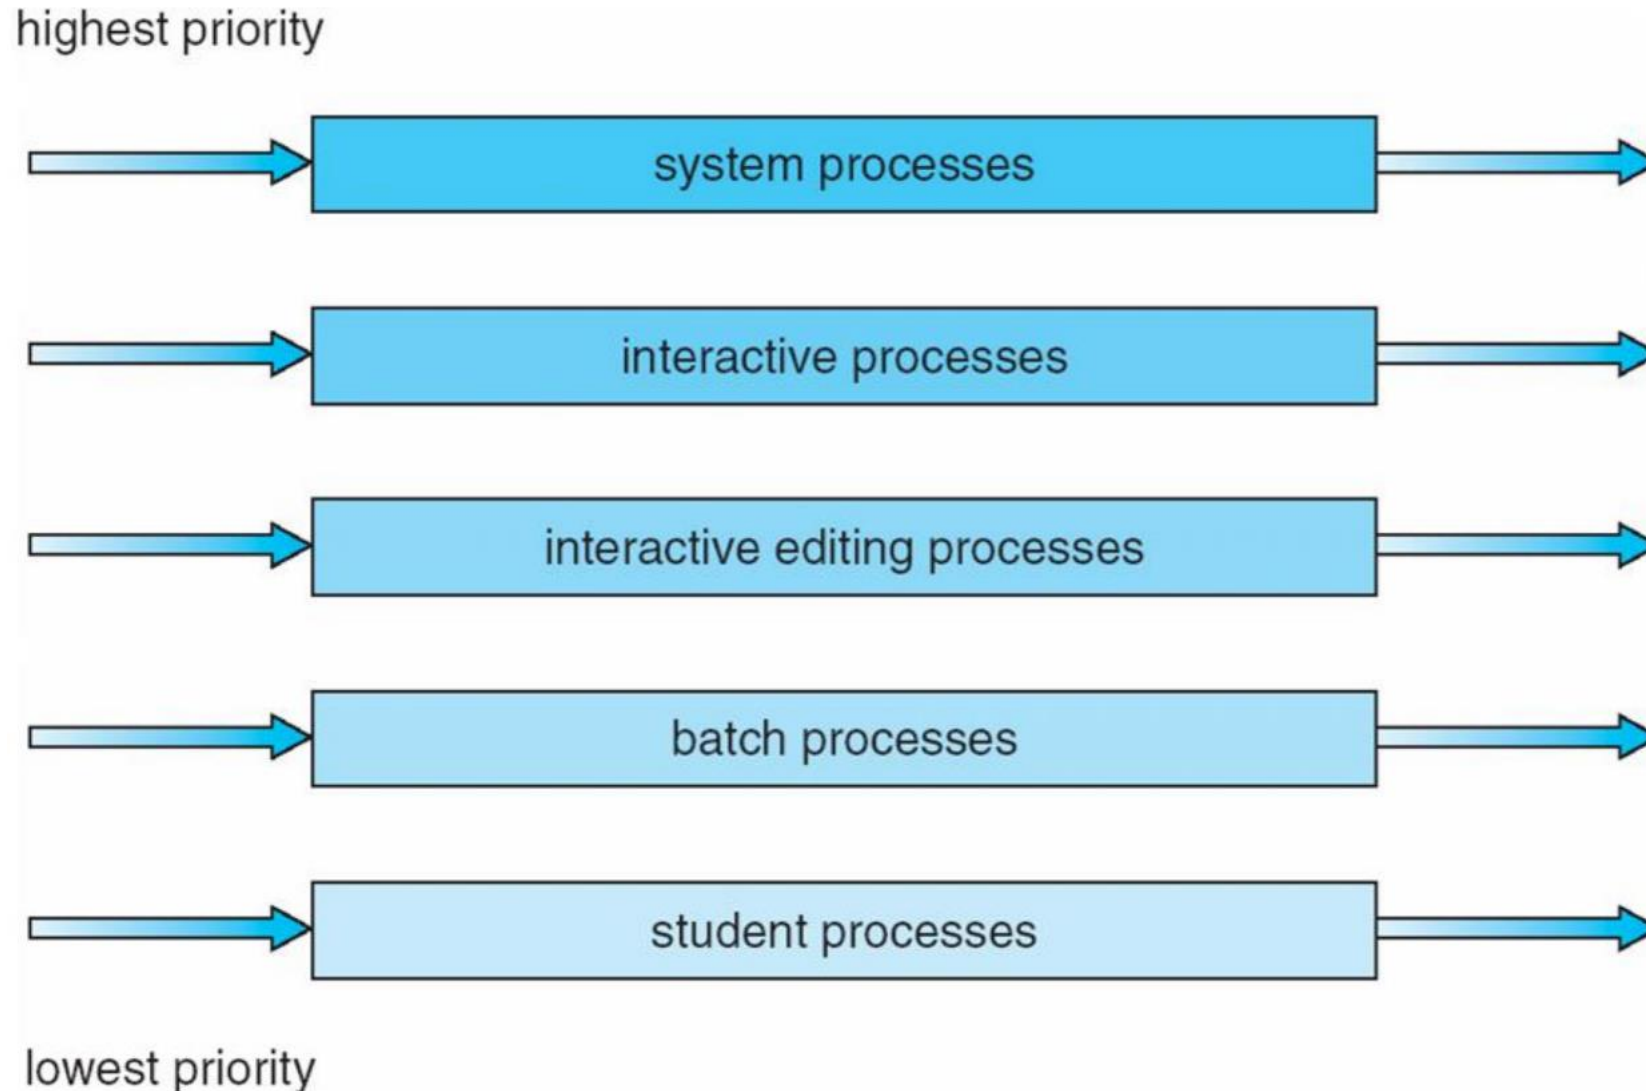
### ➤ Example:

- Run "yes &" create a yes background process. You cannot stop it by "ctrl+z".
- Enter "fg" to bring yes process into foreground. Then, you can stop it by "ctrl+z".

# Multilevel Queues

- ❑ Ready queue is partitioned into **separate queues**
- ❑ Each queue has its **own scheduling algorithm**
  - Typical example
    - ✓ foreground queue- RR
    - ✓ background queue- FCFS
- ❑ **Scheduling** must be done among these queues:
  - Fixed priority scheduling; (i.e., serve all from foreground, then from background). Possibility of starvation.
  - Time slice scheduling- each queue gets a certain amount of CPU time, e.g., 80% to foreground processes in RR, 20% to background processes in FCFS.

# Multilevel Queues



# Feedback-enabled Multilevel Queues

## ❑ Multilevel queues + **feedback**

- Process can move among queues based on their feedbacks (e.g., waiting time or service time).

## ❑ In the feedback-enabled multilevel scheduling method

- # of queues and their related scheduling strategies (i.e., **intra-queue scheduling**).
- Scheduling strategies among different queues (i.e., **inter-queue scheduling**).
- Method used to adjust the processes among queues (**feedback control**).
- Initially process placement.



# Feedback-enabled Multilevel Queues

## □ An example of feedback-enabled multilevel scheduling method

### ➤ There are three queues, and their **intra-queue scheduling** are:


- ✓ Q0 - RR with time quantum 8 ms
- ✓ Q1 - RR with time quantum 16 ms
- ✓ Q2 - FCFS

### ➤ The **inter-queue scheduling** is

- ✓ Fix priority scheduling: Q0—high, Q1—medium, Q2—low.

### ➤ The **initial process placement** and **feedback control** are designed as follows:

- ✓ A new ready process first enters Q0
- ✓ If a process does not stop/block in 8 ms, the process is moved to Q1; otherwise, it stays in Q0.
- ✓ Once the process moves to Q1, it will receive 16 ms in the next cycle.
- ✓ If the process still does not stop/block in 16 ms, it is moved to Q2.



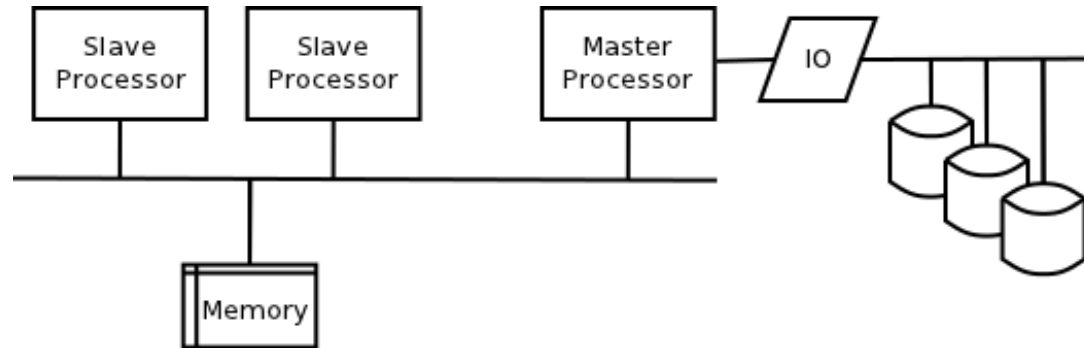
### ➤ I/O-intensive process will normally end up on high priority queue (Q0), and computational-intensive process will normally end up on low priority (Q2).

# Multiprocessor Scheduling

- ❑ What is multiprocessor scheduling
  - Given a set of runnable processes/threads, and a set of CPUs, assign processes/threads to CPUs
- ❑ Same metrics as uniprocessor scheduling
  - Fairness, efficiency, throughput, response time...
- ❑ But also new considerations
  - Load balancing
  - **Processor affinity**—keep a process running on the same core

# Asymmetric Multiprocessor Processing (AMP)

## ❑ Solution 1: Asymmetric multiprocessor processing (Centralized processing)

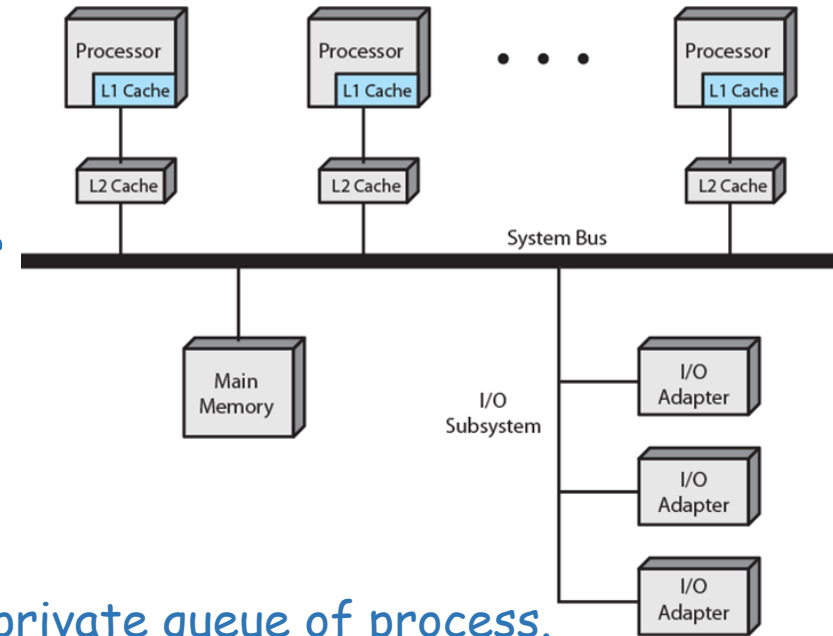


- Two types of processors (one master processor and a number of slave processors). **One ready queue(s).**
- Master processor makes scheduling decision (to balance the workload among slave processors) and handles I/O requests.
- Slave processors simply execute assigned processes.
- If a master processor fails, a slave processor become the master processor. If a slave processor fails, its allocated processes are switched to other slave processors.

# Symmetric Multiprocessor Processing (SMP)

## □ Solution 2: Symmetric multiprocessor processing (Distributed processing)

- Each processor is self-scheduling.
- All processes may be in a common ready queue (global ready queue) or each processor may have its own private queue for ready processes.
- Balancing the workload among processors is necessary to maximize the performance of the system.
  - ✓ Load balancing is necessary only if each processor has its own private queue of process.
  - ✓ Two general approaches to achieve load balancing: **Push migration** and **Pull migration**.
  - ✓ Push migration: A surveillance task periodically checks the workload on each processor and moves processes from processors with high load to processors with low load if needed.
  - ✓ Pull migration: A processor's scheduler notices its queue is empty (or less than a predefined threshold), and tries to fetch a process from another processor's queue.



# Processor affinity in SMP

## ❑ In the multiprocessor architecture

- Processors share main memory.
- Processors have their own local cache memories.
- Recently accessed data are stored in local cache memories in order to speed up data retrieval.

## ❑ Process affinity

- Try to keep the existing processes running in the same processor.
- Benefit of process affinity: quicker to restart process on same processor since the cache may already contain needed data.
- Two types of methods to achieve processor affinity:
  - ✓ **Soft Affinity**: a scheduler has a policy of attempting to keep a process running on the same processor but not guaranteeing it will do so.
  - ✓ **Hard Affinity**: some systems such as Linux have a **system call** to specify that a process shall execute on a specific processor.

```
unsigned long mask = 7; /* processors 0, 1, and 2 */
unsigned int len = sizeof(mask);
if (sched_setaffinity(0, len, &mask) < 0) {
    perror("sched_setaffinity");
}
```

# Realtime System

- ❑ In a **real-time** system, the performance of the system depends on
  - The logical result of the computation.
  - The time when the results are produced.
- ❑ **Deadline** is associated a particular task
  - **hard deadline**: required to complete a critical task within a guaranteed amount of time
  - **soft deadline**: the deadline is desirable but not mandatory; still make sense to schedule and complete the task even if it has passed its deadline.
- ❑ Two types of processes in real-time system, i.e., Periodic and aperiodic process
  - Periodic processes: arriving at fixed frequency, can be characterized by 3 parameters (C,D,T) where C = service/burst time, D = relative deadline, T = period (e.g., 20ms, or 50HZ). Periodic processes are called Time-driven processes, their activations are generated by timers.
  - Aperiodic processes : all processes that are not periodic, also known as event-driven, their activations may be generated by external interrupts.

# Realtime Scheduling


## □ Realtime scheduling

- Need **preemptive** strategy & priority function
- The service time of each process is known in advance.
- Find a schedule for all the processes so that each meets its deadline.
- A popular algorithm: EDF (**Earliest Deadline First**)

# Earliest Deadline First (EDF) Scheduling

□ Scheduler selects a job (e.g., a process) with EDF

- The highest priority job is the one with the earliest **absolute deadline**;
  - ✓ Absolute deadline vs relative deadline




- 
- If two jobs (one is the executed by the CPU and the other is in the ready queue) have the same absolute deadline, select the running job.
- Decision mode: **preemption**.



# Earliest Deadline First (EDF) Scheduling

## □ Example: scheduling with EDF

- ✓ Three types of processes  $\tau_1, \tau_2, \tau_3$  are initially in the ready queue.
- ✓  $\tau_1=(1,4,4)$ , which indicates that the service time of  $\tau_1$  is 1 time unit, and  $\tau_1$  will be in the ready queue after each 4 units; the relative deadline of  $\tau_1$  is also 4 units. Accordingly,  $\tau_2=(2,6,6)$ , and  $\tau_3=(3,8,8)$ .



$$CPU\ utilization = \frac{23}{24}$$

# Earliest Deadline First (EDF) Scheduling

- **Theorem:** Given a set of periodic or sporadic jobs, with relative deadlines equal to periods, the job set can be schedulable by EDF iff

$$U = \sum_{i=1}^N \frac{C_i}{T_i} \leq 1$$

where  $C_i$  is the service time (burst time) of job  $i$ ,  $T_i$  is the relative deadline of job  $i$ ,  $N$  is the total number of jobs in the job set, and  $U$  is the CPU utilization.


- **Lemma:** EDF is an optimal algorithm, in the sense that if a job set is schedulable, then it is always schedulable by EDF.

- ✓ If  $U > 1$ , no algorithm can successfully schedule the job set;
- ✓ If  $U \leq 1$ , EDF can always provide a feasible schedule.

# Earliest Deadline First (EDF) Scheduling

## □ Domino effect with EDF


- If  $U > 1$  (i.e., the job is NOT schedulable), we have the **domino effect** with EDF: it means that many jobs miss their deadlines.
- An example of domino effect:
  - ✓ Four processes:  $\tau_1$ ,  $\tau_2$ ,  $\tau_3$ , and  $\tau_4$ .
  - ✓  $\tau_1 = (2, 5, 5)$ ,  $\tau_2 = (2, 6, 6)$ ,  $\tau_3 = (2, 7, 7)$ ,  $\tau_4 = (2, 8, 8)$



# Rate Monotonic(RM) Scheduling

## □ Rate Monotonic (RM) Scheduling

- The highest priority job is the one with the earliest **relative deadline**;
- Decision mode: preemption (by default).
- Four processes:  $\tau_1$ ,  $\tau_2$ ,  $\tau_3$ , and  $\tau_4$ , where  $\tau_1 = (2, 5, 5)$ ,  $\tau_2 = (2, 6, 6)$ ,  $\tau_3 = (2, 7, 7)$ ,  $\tau_4 = (2, 8, 8)$ 
  - ✓  $\tau_1$  and  $\tau_2$  never miss their deadlines;
  - ✓  $\tau_3$  misses many deadlines;
  - ✓  $\tau_4$  is not executed!



# Rate Monotonic(RM) Scheduling

## □ Rate Monotonic (RM) Scheduling

- RM cannot guarantee all the jobs meet their deadlines, even if the job set is schedulable (i.e.,  $U \leq 1$ ).
- $\tau_1=(1,4,4)$ ,  $\tau_2=(2,6,6)$ , and  $\tau_3=(3,8,8)$ .  $U = \frac{1}{4} + \frac{2}{6} + \frac{3}{8} = \frac{23}{24}$

