

ECE437/CS481

# M04A: PROCESS COORDINATION RACE CONDITION & CRITICAL SECTION

CHAPTER 5.1-5.4

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a small upward curve on the left, dipping into a V-shape in the center, and then curving back up on the right before continuing as a straight line to the edge.

# Dependent & Independent Processes

- ❑ Independent processes: **not affected by rest of universe**
  - No shared resource (memory, filesystem) among processes
  - Deterministic: results depends on the inputs
  - Block and restart without adverse effects
- ❑ Dependent/Cooperative processes: **Non-independent**
  - shared some resources with other processes
  - Non-deterministic: may have different results each time
  - May be irreproducible: difficult to debug and test

# Cooperating Requires Synchronization

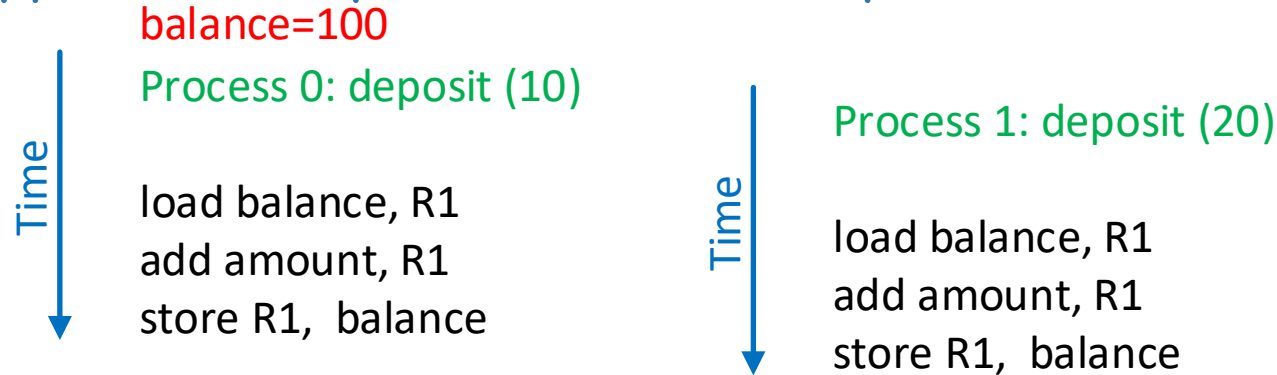
## ❑ Two processes share **account balance**

```
void deposit (int amount) {  
    balance += amount;  
}
```

### ➤ “Making deposit” can be complied into

```
load balance, R1  
add amount, R1  
store R1, balance
```

### ➤ What happens if two processes make deposit at same time?



### ➤ What could be the final balance? 100, 110, 120, or 130?

### ➤ **Race condition: when two or more processes/threads are accessing the shared resource, the result depends on the ordering of interleaving.**

# Cooperating Requires Synchronization

## ❑ Deal with race condition

➤ To avoid race condition → setup a **critical section (CS)**

- ✓ A **critical section** is a **code segment** intended to access a shared resource (file, input or output port, global variable, etc.) that must NOT be concurrently accessed by more than one process/thread of execution.

```
void deposit (int amount) {  
    disable accessing balance by other processes;  
    balance += amount;  
    enable accessing balance by other processes;  
}
```

# Cooperating Requires Synchronization

## □ Properties/Requirements of a critical section

### ➤ Mutual exclusion

- ✓ Only one process/thread in critical section at a time.

### ➤ Guaranteed progress

- ✓ If no process/thread is in its critical section, and if one or more processes/threads are waiting to execute the critical section, then any one of these processes/threads must be allowed to get into the critical section.

### ➤ Bounded waiting

- ✓ Must eventually allow waiting processes/threads to proceed.

### ➤ Solution should be independent on

- ✓ # of CPUs, # of processes/threads, and speed of CPUs.

# Cooperating Requires Synchronization

## □ How to establish/realize a CS?

- Disable interrupts, which essentially disable the context switch.

```
disable interrupt;  
Critical Section (CS);  
enable interrupt;
```

- ✓ Not work if there are multiple CPUs. Disabling interrupts affects just one CPU.
- ✓ Unsafe to give the system privilege to user.

# Cooperating Requires Synchronization

## □ How to establish/realize a CS?

### ➤ Using lock variable

```
shared bool lock = FALSE;
```

```
Process A
```

```
while (lock == TRUE)
```

```
    /* null */ ;
```

```
lock = TRUE;
```

```
    CS;
```

```
lock = FALSE;
```

```
Process B
```

```
while (lock == TRUE)
```

```
    /* null */ ;
```

```
lock = TRUE;
```

```
    CS;
```

```
lock = FALSE;
```

- ✓ Not work since setting the lock does not satisfy mutual exclusion, thus the CS is not mutual exclusion.

# Cooperating Requires Synchronization

## □ How to establish/realize a CS?

### ➤ Using separated lock variables

```
shared bool lock0 = FALSE, lock1 = FALSE;
```

Process A

```
lock0 = TRUE;
```

```
while (lock1 == TRUE)
```

```
    /* null */ ;
```

```
CS;
```

```
lock0 = FALSE;
```

Process B

```
lock1 = TRUE;
```

```
while (lock0 == TRUE)
```

```
    /* null */ ;
```

```
CS;
```

```
lock1 = FALSE;
```

- ✓ Can achieve mutual exclusion in single core system.
- ✓ Bounded waiting does not satisfy, i.e., the two processes are waiting for each other forever.



# Cooperating Requires Synchronization

## □ How to establish/realize a CS?

### ➤ Using turn variable

shared int turn = 0,

#### **Process A**

while (turn != 0)

/\* null \*/ ;

CS;

turn = 1;

#### **Process B**

while (turn != 1)

/\* null \*/ ;

CS;

turn = 0;

- ✓ Can achieve mutual exclusion in single core system.
- ✓ Guaranteed progress does not satisfy.

# Cooperating Requires Synchronization

## □ How to establish/realize a CS?

### ➤ Using turn + separated locks (Peterson's solution)

	<code>shared int turn=0; shared bool lock[2] = {FALSE, FALSE};</code>	
	<code>Process A</code>	<code>Process B</code>
I request to access the CS	→ <code>lock[0] = TRUE;</code>	<code>lock[1] = TRUE; /*req*/</code>
Assume that it is your turn	→ <code>turn = 1;</code>	<code>turn = 0;</code>
If you request to access the CS and it is your turn, then I will wait.	→ <code>while (lock[1] &amp;&amp; turn==1)</code>	<code>while (lock[0] &amp;&amp; turn==0)</code>
	<code>/* null */ ;</code>	<code>/* null */ ;</code>
	<b>CS;</b>	<b>CS;</b>
	<code>lock[0] = FALSE;</code>	<code>lock[1] = FALSE;</code>

- ✓ Can achieve mutual exclusion in single core system. But it is still unable to guarantee progress.
- ✓ Complicated and error-prone when dealing with N processes.
- ✓ A simple and low-level solution is needed.