

ECE437/CS481

# M06B: DISK SPACE MANAGEMENT

CHAPTER 12.4-12.5


Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a thin line, dipping into a V-shape, and then rising back to a thin line, creating a stylized horizon or wave effect.

# File Space Allocation

## □ Disk Structure



- **Sector:** Minimum unit (normally, 512 Byte for old hard disks; 4 KB for new hard disks) for data storage.
  - ✓ Use `cat /sys/block/sda/queue/physical_block_size` to see the size of sector for your hard disk.
- **Track:** One cycle on the platter
- **Cylinder:** The tracks with the same track number on different platters.
- Each sector is identified by a combination of cylinder number, the head number, and the sector number.

# File Space Allocation

## ❑ Logic blocks in file systems

- The OS uses logic **blocks** to read and write files.
  - ❑ Size of a block =  $2^n \times$  size of a sector
  - ❑ "stat -f /" to check the size of a block
  - ❑ The way of addressing logic blocks is a simple linear addressing scheme
  - ❑ Mapping a logic block to a sector(s) is required

LBA	C	H	S
0	0	0	0
1	0	0	1
2	0	0	2
3	0	0	3
4	0	0	4
5	0	0	5
6	0	0	6
7	0	0	7
8	0	0	8
9	0	0	9
10	0	1	0
11	0	1	1
12	0	1	2
13	0	1	3
14	0	1	4
15	0	1	5
16	0	1	6
17	0	1	7
18	0	1	8
19	0	1	9

Cylinder 0

LBA	C	H	S
20	1	0	0
21	1	0	1
22	1	0	2
23	1	0	3
24	1	0	4
25	1	0	5
26	1	0	6
27	1	0	7
28	1	0	8
29	1	0	9
30	1	1	0
31	1	1	1
32	1	1	2
33	1	1	3
34	1	1	4
35	1	1	5
36	1	1	6
37	1	1	7
38	1	1	8
39	1	1	9

Cylinder 1

$$\text{LBA} = ((C \times \text{HPC}) + H) \times \text{SPT} + S$$

- ❖ C, H and S are the cylinder number, the head number, and the sector number
- ❖ LBA is the logical block address
- ❖ HPC is the number of heads per cylinder
- ❖ SPT is the number of sectors per track

# File Space Allocation


- ❑ How to allocate blocks to these files so that disk space is utilized effectively, and files can be accessed quickly?—File Space Allocation
- ❑ Three major methods of allocating disk blocks:
  - Contiguous allocation
  - Linked allocation
  - Indexed allocation

# File Space Allocation

## ❑ Continuous allocation

- Continuous blocks is allocated to a file.
- FCB only needs to store the information of **the 1st block** and the **# of blocks allocated**.
- Good performance for **sequential access**
- Easy to achieve **direct access**
  - ✓ If the system tries to access block  $i$  of a file that starts at block  $b$ , the system can immediately access block  $b + i$

name	1 <sup>st</sup> block	length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3



# File Space Allocation

## Continuous allocation

- Block allocation problem: given the file size, how to pick up contiguous group of blocks to store the file.


### First fit:

- choose the first unused contiguous group of blocks of sufficient sizes.

### Best fit:

- choose the smallest unused contiguous group of blocks of sufficient sizes.

- If the system tries to allocation a file with size equal to **three logic blocks**, which logic blocks will be allocated to the file by applying FF and BF, respectively?



# File Space Allocation


## ❑ Drawbacks of continuous allocation

### ➤ Suffers from the problem of **external fragmentation**

- ✓ External fragmentation: As files are allocated and deleted, the free disk space is broken into little pieces. Thus, none of these pieces is large enough to store files.
- ✓ Solution: Copy an entire file system onto another disk, and then copy the file system back to the original disk by allocating contiguous space.

### ➤ May have problems in estimating the size of a file.


- ✓ Determining how much space is needed for a file in continuous allocation.
- ✓ In general, estimating the size of a file is difficult.
- ✓ If the estimated size of a file is too small, the file may not be feasible to be extended (especially for the best-fit allocation strategy).



# File Space Allocation

## □ Linked allocation

- Each file is a **linked list of blocks** which need NOT be contiguous. The blocks can be scattered anywhere on the disk.
  - ✓ A file of five blocks might start at block 9 and continue at block 16, then block 1, then block 10, and finally block 25.
- Each block contains a pointer to the next
  - ✓ These pointers are not made available to the OS.
- FCB only contains **the 1st and last blocks of the file**





# File Space Allocation

## □ Linked allocation pros and cons

### ➤ Flexible to extend blocks of a file

- ✓ A write to the file causes to find a free block, and this new block is written to and is linked to the end of the file.
- ✓ To read a file, we simply read blocks by following the pointers from block to block.
- ✓ The size of a file need not be declared when the file is created. A file can continue to grow as long as free blocks are available.

### ➤ No external fragmentation

- ✓ Any free block can be used to satisfy a request.

### ➤ Ineffectively for direct access

- ✓ To find the  $i^{\text{th}}$  block of a file, we must start at the beginning of that file and follow the pointers until we get to the  $i^{\text{th}}$  block.

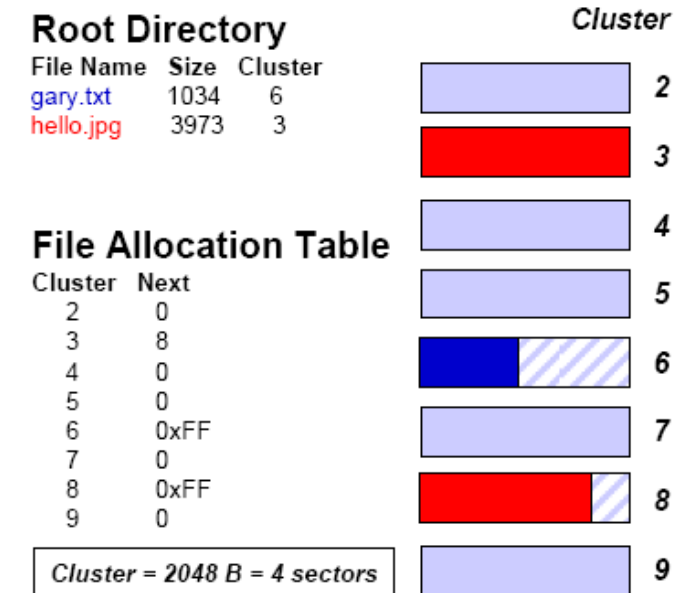
### ➤ Extra overhead for storing points in blocks

### ➤ Low reliability

# File Space Allocation

## ❑ Linked allocation: MS-DOS FAT as an example

- FAT = File Allocation Table, a variation of linked allocation
  - ✓ one FAT per file system
  - ✓ one entry per block containing the next block in the file
- Normally, one block = 512 bytes → too small
- Use cluster in FAT to replace block, where one cluster =  $2^x$  blocks ( $x=0,1,2,3,4,5,6,7\dots$ )
- Size of a FAT is decided by total # of clusters on storage
  - ✓ In FAT, each entry represents a cluster
  - ✓ "0" means the cluster is free
  - ✓ "FF" means EOF



# File Space Allocation

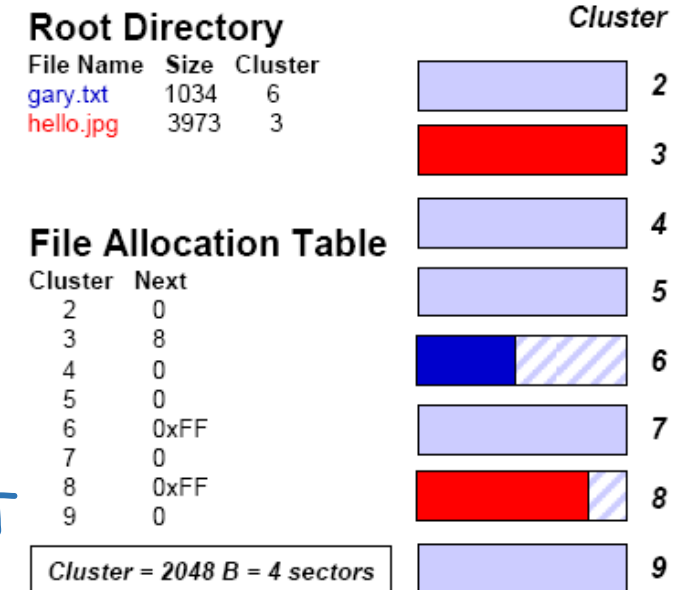
## □ Linked allocation: MS-DOS FAT as an example

### ➤ FAT-12: 12-bit addr space to index the entries in FAT

- ✓ There are maximum  $2^{12}$  clusters
- ✓ If one cluster=one block= 512 B, then  
max volume size=512 B/cluster\*  $2^{12}$  clusters=2 MB---floppy disk,  
whose size is 1.44 MB
- ✓ The size of a cluster (i.e.,  $2 \times$  block) is specified in the boot record.

### ➤ FAT-16 (1987): 16-bit addr space to index the entries in FAT

- ✓ There are maximum  $2^{16}$  clusters
- ✓ If one cluster=one block= 512 B, then max volume size=512 B/cluster\*  $2^{16}$  clusters=32 MB  
**Too small**
- ✓ If one cluster=128 blocks= 64 KB, then max volume size= 64 KB/cluster \*  $2^{16}$  clusters=4 GB.  
Fairly enough but incurs **internal fragmentation**.
- ✓ **Internal fragmentation** is the wasted space within each allocated block because of rounding up from the actual requested allocation to the allocation granularity.



# File Space Allocation

## □ Linked allocation: MS-DOS FAT as an example

➤ FAT-32 (1996, Windows 95): 32-bit addr space to index the entries in FAT

- ✓ There are maximum  $2^{32}$  clusters
- ✓ If one cluster=one block= 512 B, then max volume size= $512 \text{ B}/\text{cluster} * 2^{32} \text{ clusters}=2 \text{ TB}$
- ✓ If one cluster=8 blocks= 4 KB, then max volume size= $4 \text{ KB}/\text{cluster} * 2^{32} \text{ clusters}=16 \text{ TB}$


The setup is used as the default setting by NTFS

➤ FAT requires storing the entire table in memory for efficient access --- A huge table!

# File Space Allocation

## Indexed allocation

- Bring all pointers together in an index block
- The  $i^{\text{th}}$  entry in the index block points to the  $i^{\text{th}}$  block of the file.
- FCB contains a pointer to its index block.



## □ Indexed allocation pros and cons

- Flexible to extend, insert, and delete blocks
- Index allocation supports both sequential and direct access without external fragmentation.
- Suffer from wasted space
  - ✓ The overhead of the index block may be greater than the pointer overhead of linked allocation.
- An index block could be too small to hold enough pointers for a large file

# File Space Allocation

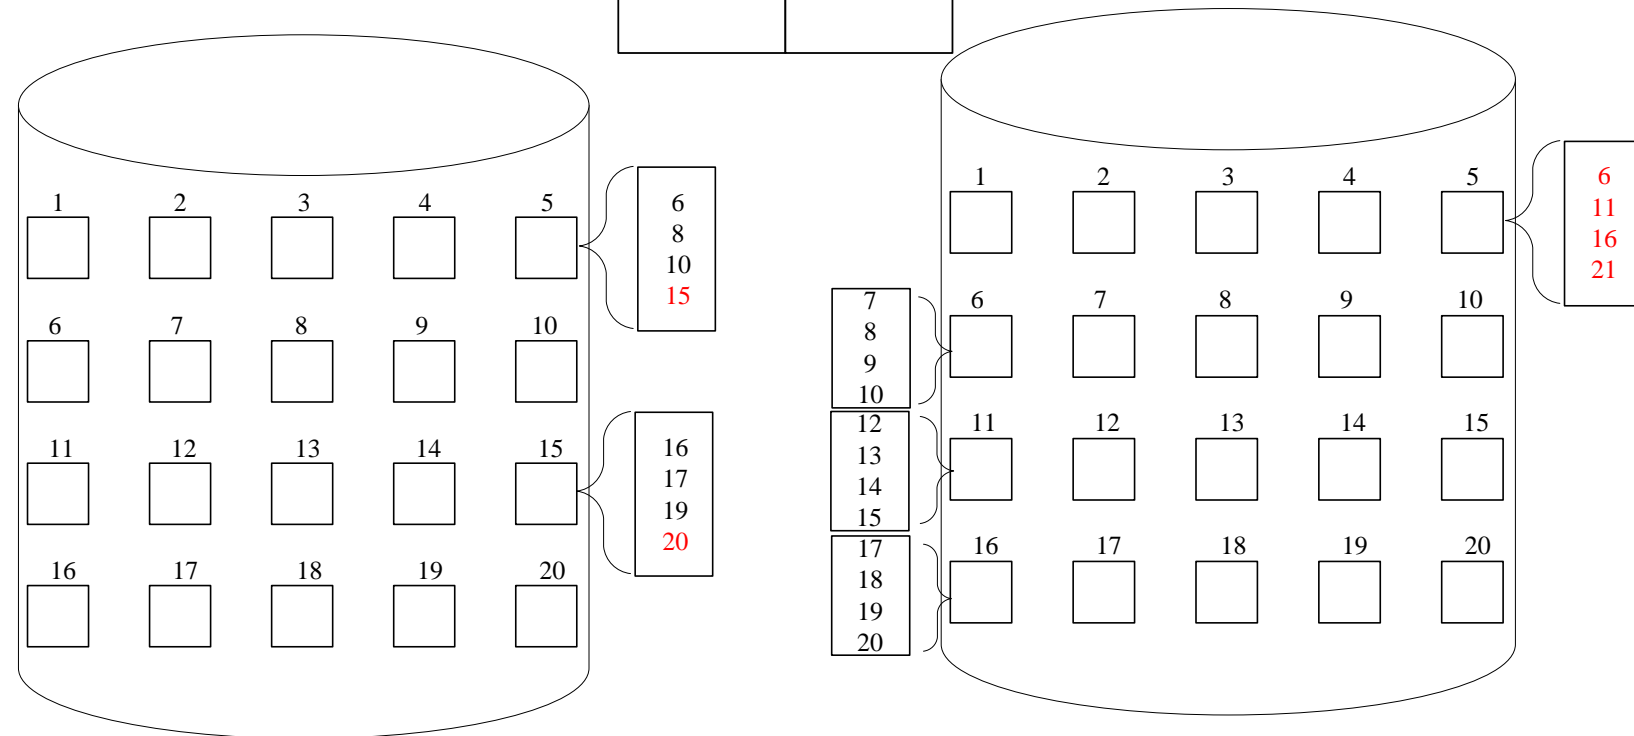
## ❑ Indexed allocation (Multiple index blocks in UNIX-based file system)

### ➤ Solution

- 1) Have multiple index blocks and chain them into a linked-list.
- 2) Have multiple index blocks but make them a tree just like the indexed access method.

Directory Table

File Name	Index block
Example.txt	5

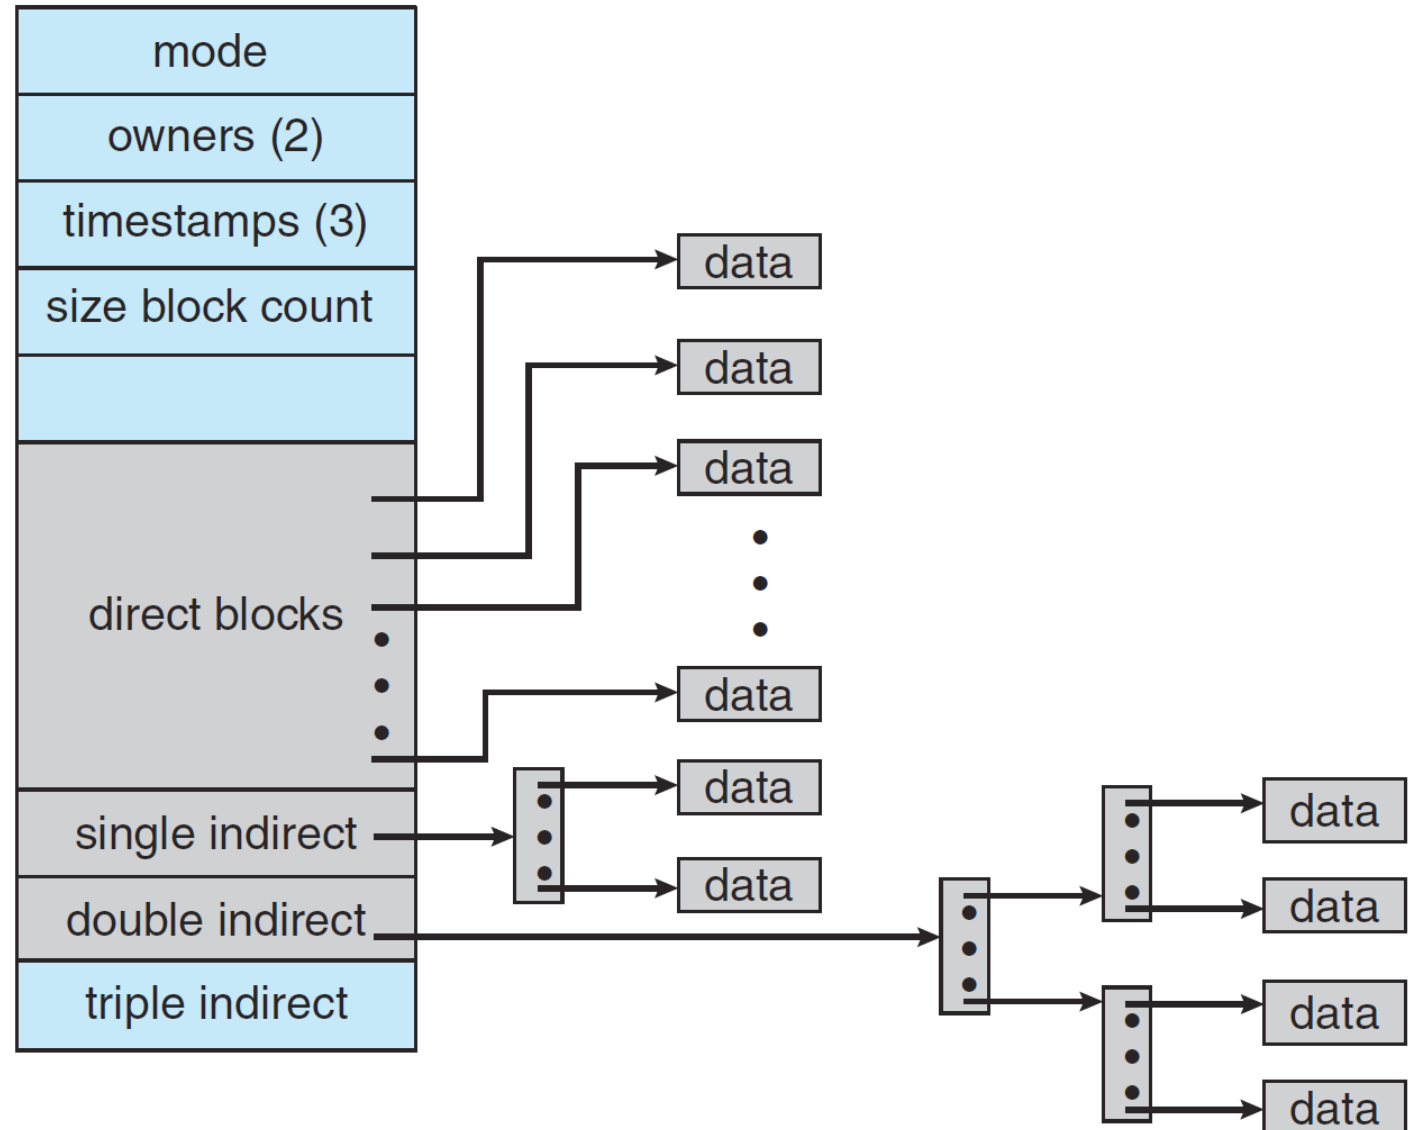


# File Space Allocation

## □ Indexed allocation (Multiple index blocks in UNIX-based file system)

### ➤ Solution

#### 3) Mixed solution





# Free Space Management

- ❑ How do we keep track free blocks on a disk?
- ❑ A free block list is maintained. When a new block is requested, we search this list to find one.
- ❑ The following are commonly used techniques:
  - ✓ Bit Vector
  - ✓ Linked List
  - ✓ Grouping
  - ✓ Linked List + Address + Count

# Free Space Management

## ❑ Bit Vector

- Each block is represented by a bit in a vector. Thus, if there are  $n$  disk blocks, the vector has  $n$  bits.
- If a block is free, its corresponding bit is 0.
- When a block is needed, the vector is searched.
- If the disk capacity is small, the whole bit vector can be stored in the main memory. For a large disk, this bit vector will consume too much memory.
- We could group a few blocks into a **cluster** and allocate **clusters**. This can reduce the size of the vector, but may cause internal fragmentation.



$$\text{bit}[i] = \begin{cases} 0 \Rightarrow \text{block}[i] \text{ free} \\ 1 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$


## □ Bit Vector

- Block size=512 B= $2^9$  Bytes
- Disk size=1 TB= $2^{40}$  Bytes
- What is the size of the bit vector?
  
- Solution:
  - ✓ Total number of blocks= $2^{40} / 2^9 = 2^{31}$
  - ✓  $2^{31}$  bits (=256 MB) are used to represent  $2^{31}$  blocks.
  - ✓ That is, the size of the bit vector is 256 MB, which is normally stored in the memory.
  - ✓ If we group 4 blocks into a cluster, size of the bit vector= 64 MB.

# Free Space Management

## ❑ Linked List


- Like the linked allocation method, free blocks can be chained into a linked list.
- When a free block is needed, the first one in the chain is allocated.
- However, this method has the **same disadvantages of the linked allocation method**.
- We may use a FAT for the disk and chain the free block pointers together. Note that the FAT may be very large and consume memory space.



### Root Directory

File Name	Size	Cluster
gary.txt	1034	6
hello.jpg	3973	3

### Cluster



### File Allocation Table

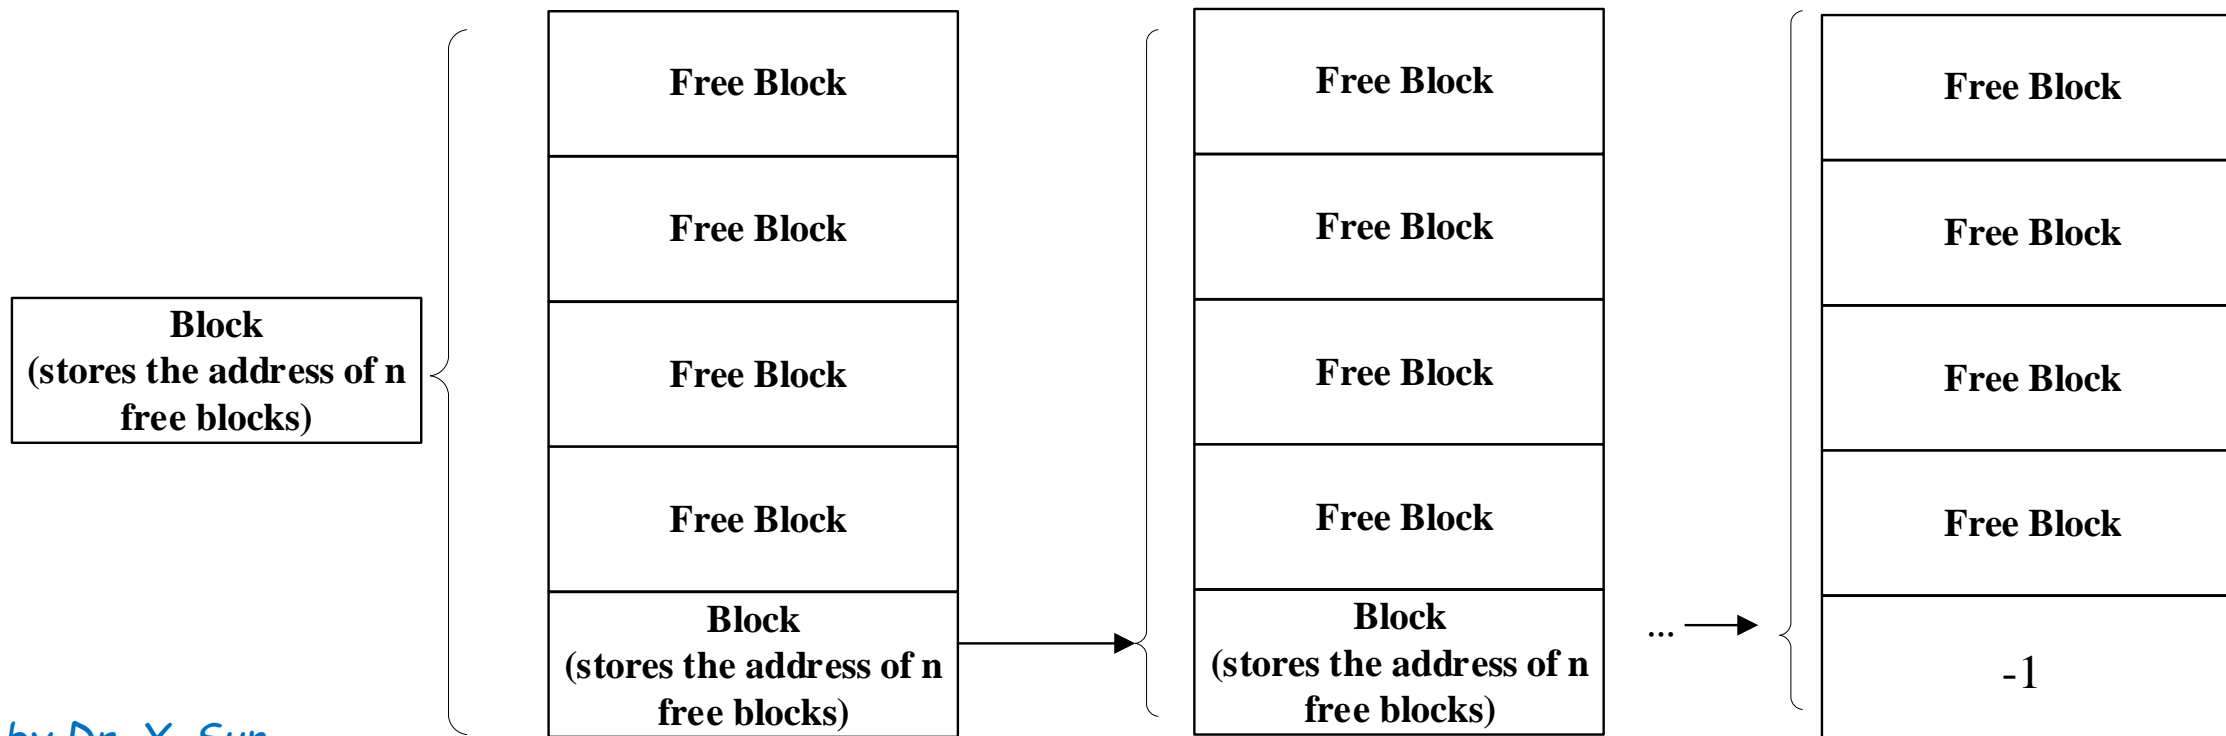
Cluster	Next
2	0
3	8
4	0
5	0
6	0xFF
7	0
8	0xFF
9	0

Cluster = 2048 B = 4 sectors

# Free Space Management

## ❑ Grouping

- The first free block stores the address of  $n$  free blocks.
- Out of these  $n$  blocks, the first  $n-1$  blocks are actually free, and the last block contains the address of next free  $n$  blocks.
- In this way, we can quickly locate free blocks.



## □ Grouping

- Block size=4 KB= $2^{12}$  Bytes
- Disk size=1 TB= $2^{40}$  Bytes
- Assume that 32 bits are used to represent a block number
- How many blocks are needed to store the addresses of free blocks in the worst-case scenario?
  
- Solution:
  - ✓ Each block can store  $2^{12}/4=2^{10}$  block numbers.
  - ✓ In the worst-case scenario, total number of free blocks=  $2^{40} / 2^{12} = 2^{28}$
  - ✓ Number of needed blocks=  $2^{28} / 2^{10} = 2^{18}$ , which needs  $2^{18} * 4 \text{ KB} = 1 \text{ GB}$ .

# Free Space Management

## ❑ Link List + Address + Counting

- Blocks are often allocated and freed in groups.
- We can store the address of the first free block and the number of the following  $n$  free blocks.

